

Data Science Project

Team nr: 16	Student 1: Antero Morgado IST nr: 1119213 Student 2: David Ferreira IST nr: 1107077 Student 3: José Fernandes IST nr: 1103727 Student 4: Olha Buts IST nr: 1116276
--------------------	---

CLASSIFICATION

1 DATA PROFILING

May be used to describe any useful observation about the data, and that was used in the current project. An example is the use of any domain knowledge to process the data or evaluate the results. **Shall not exceed 200 characters.**

Data Dimensionality

Regarding the Missing Values analysis, the completeness metric for Dataset 1 was adjusted to treat 'UNKNOWN' occurrences as missing values.

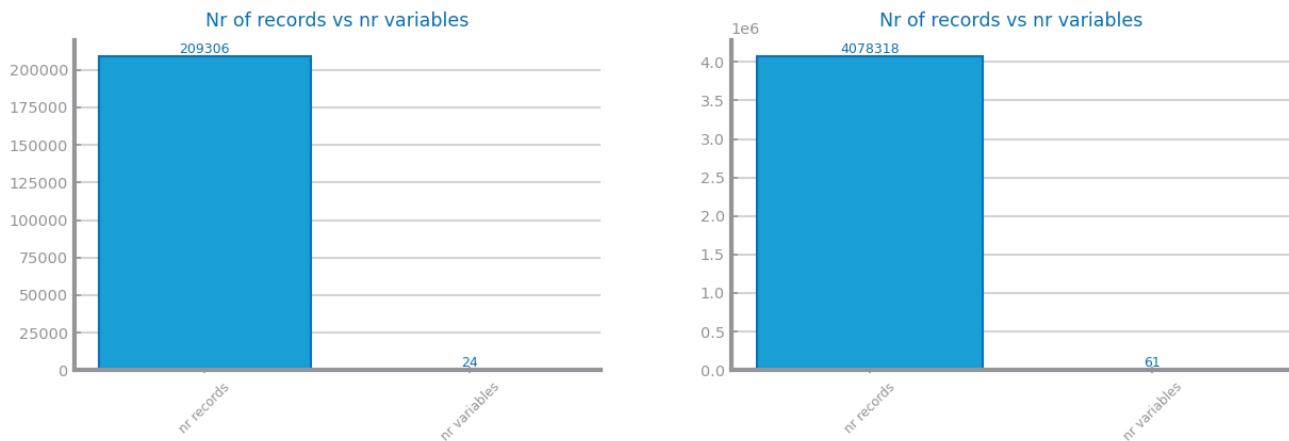


Figure 1: Nr Records x Nr variables for dataset 1 (left) and dataset 2 (right)

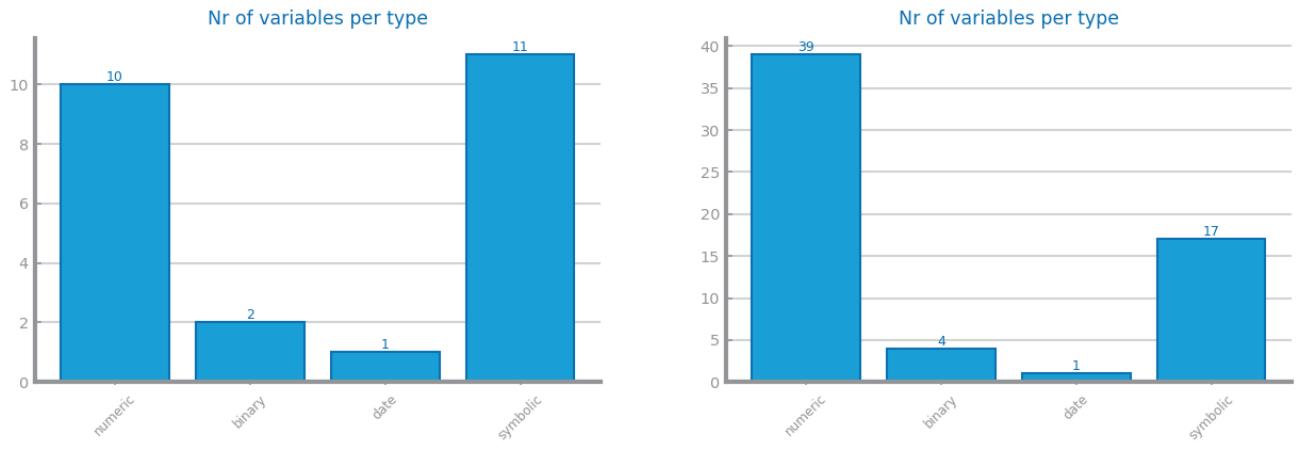


Figure 2: Nr variables per type for dataset 1 (left) and dataset 2 (right)

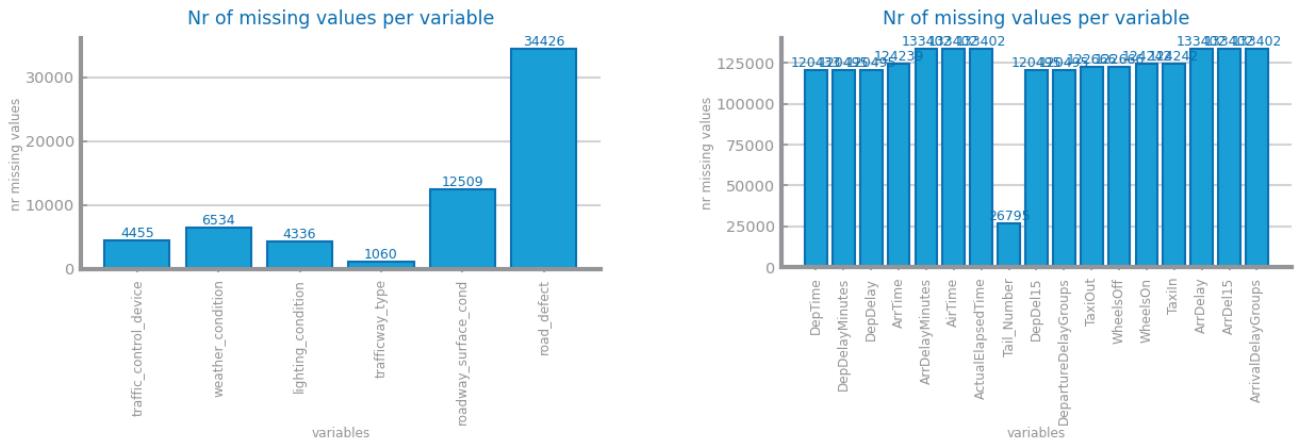


Figure 3: Nr missing values for dataset 1 (left) and dataset 2 (right)

Data Distribution

The distribution analysis reveals that the raw data is significantly dispersed, exhibiting high variance across several features. Certain variables show a clear lack of balance, with distributions heavily skewed toward specific classes or ranges and presence of evident outliers.

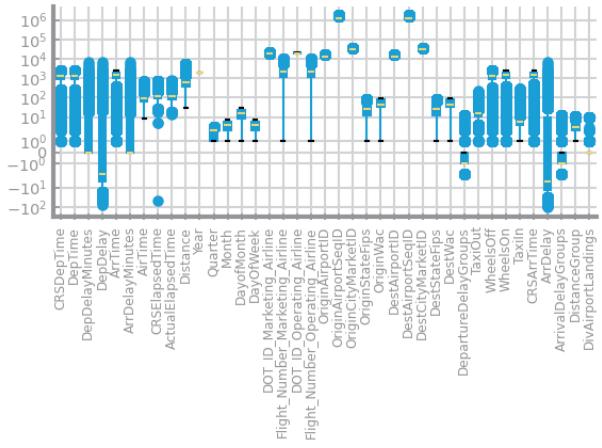
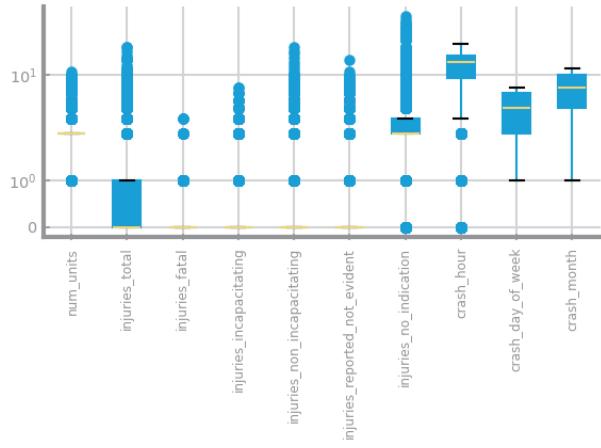


Figure 4: Global boxplots dataset 1 (left) and dataset 2 (right)

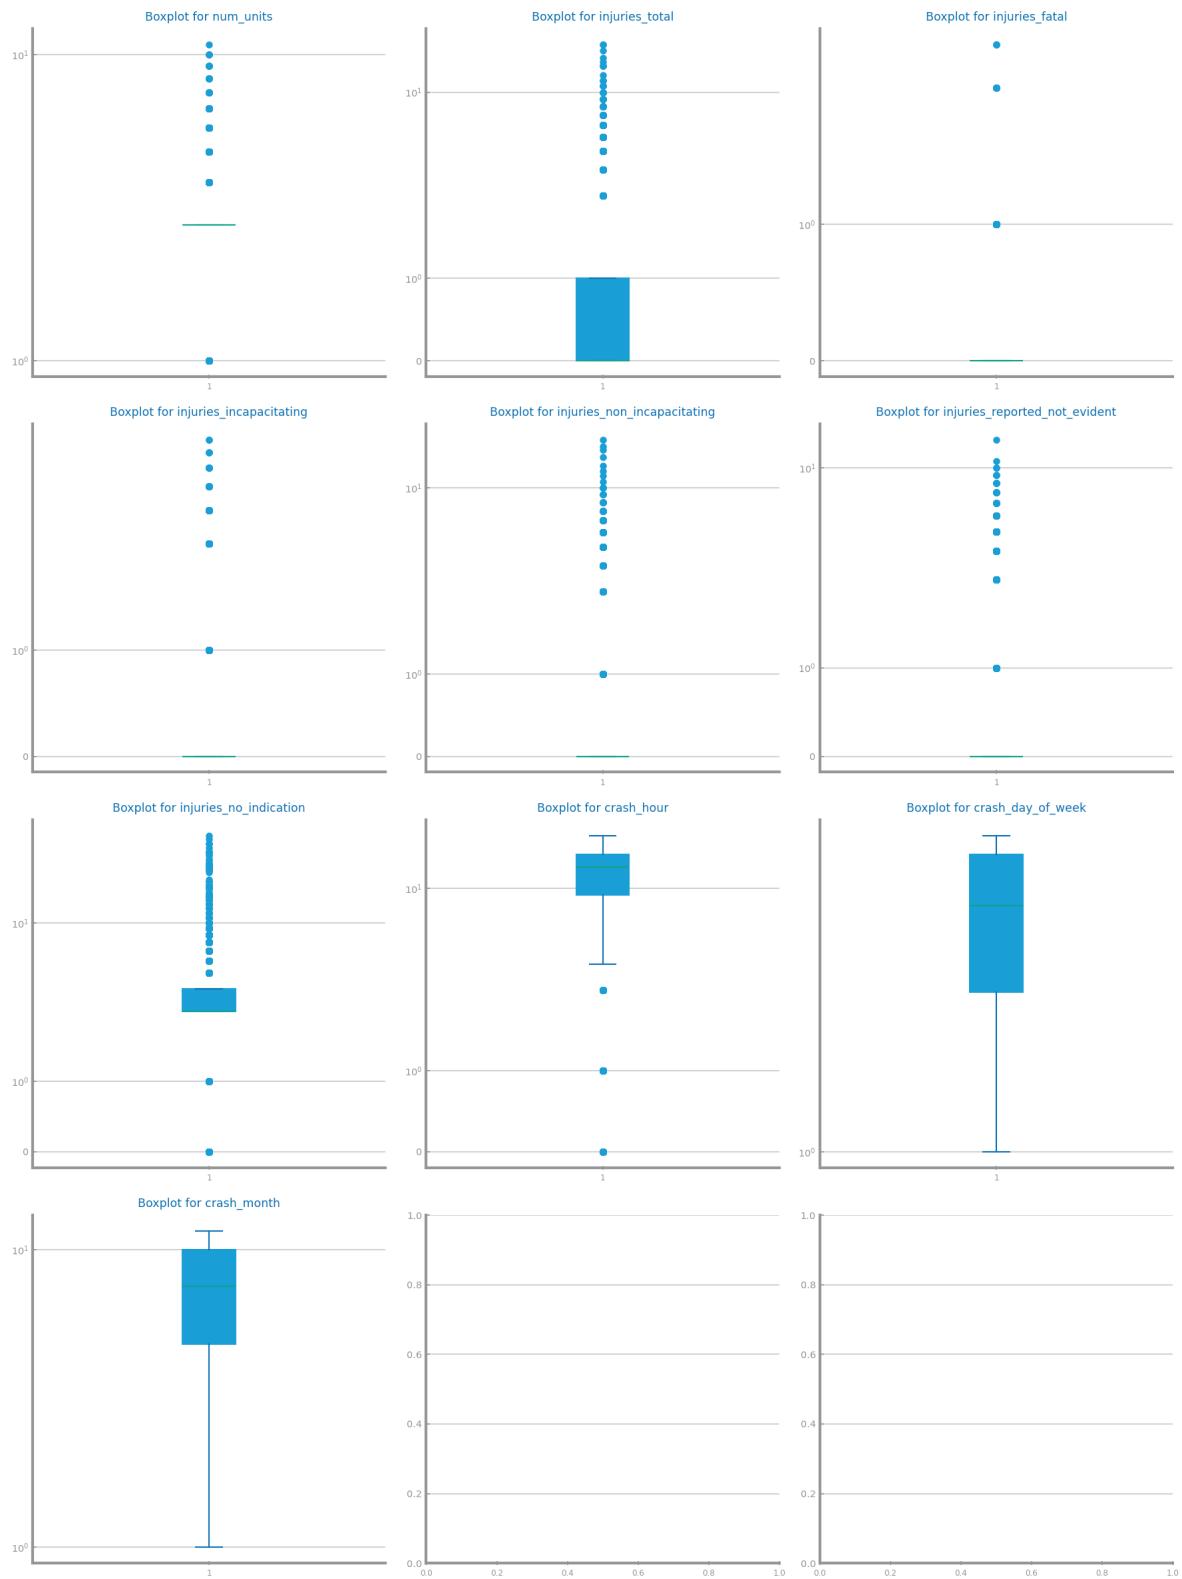


Figure 5: Single variables boxplots for dataset 1

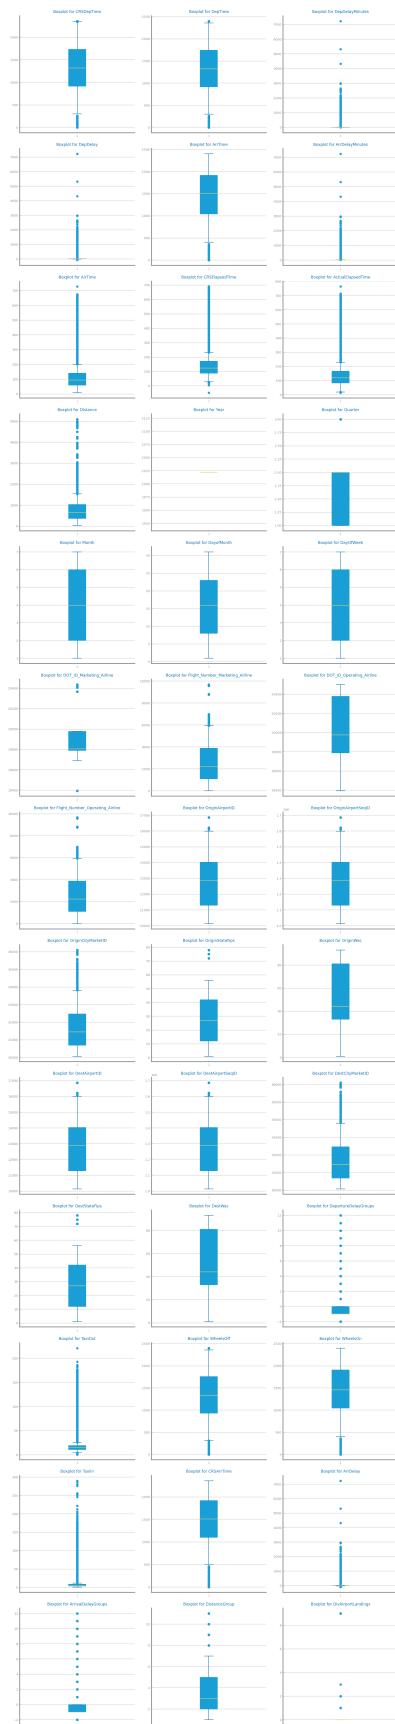


Figure 6: Single variables boxplots for dataset 2

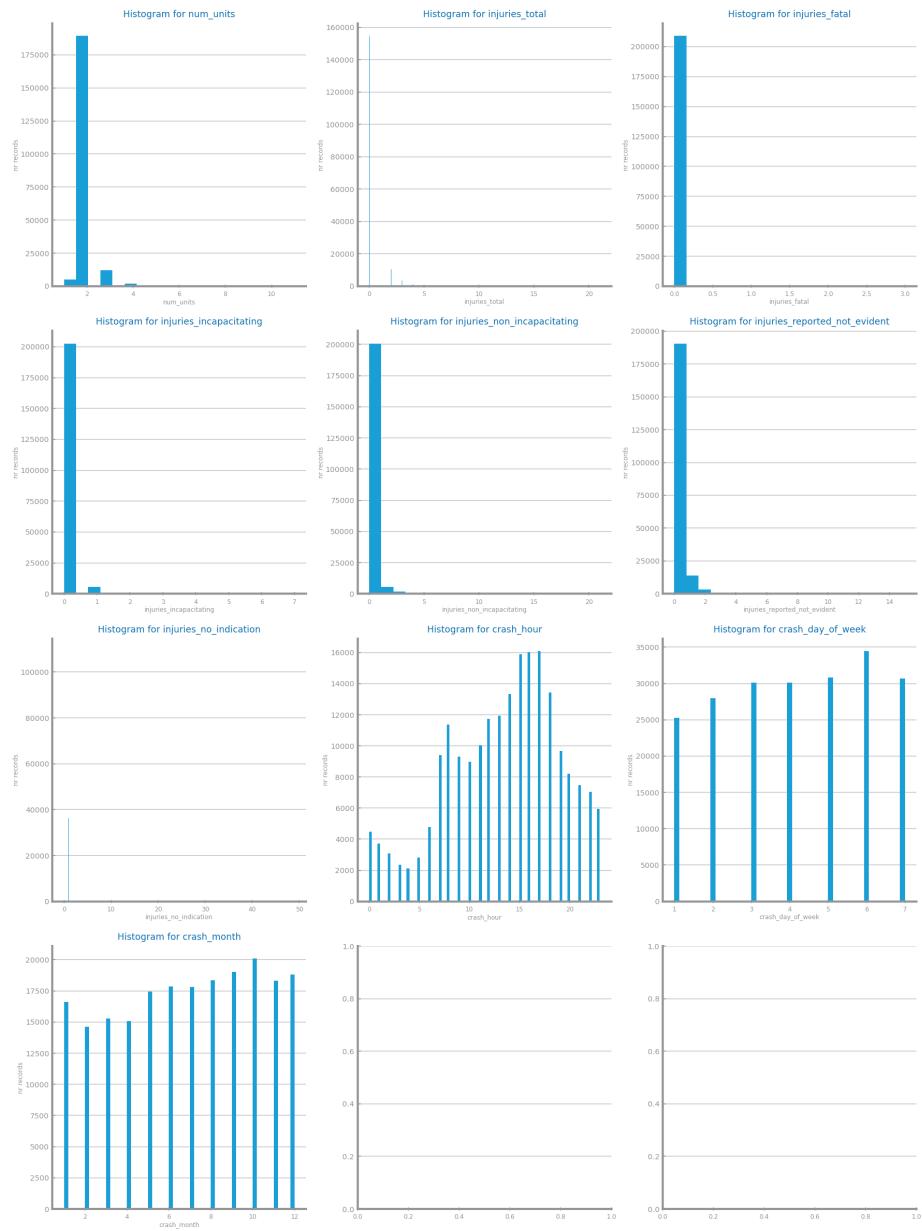


Figure 7: Histograms for dataset 1

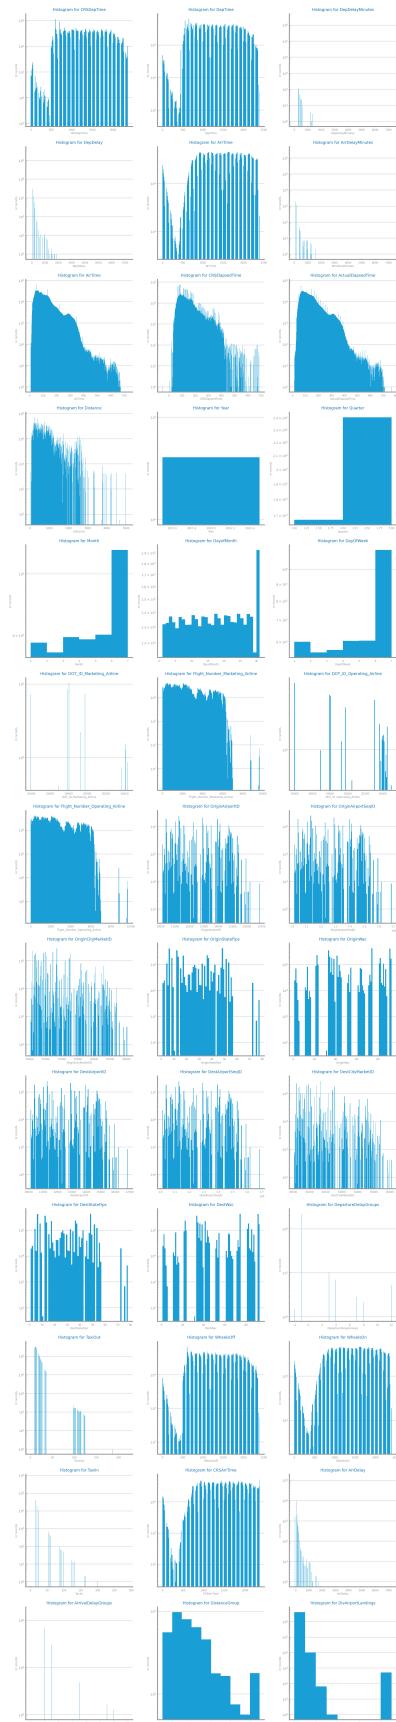


Figure 8: Histograms for dataset 2

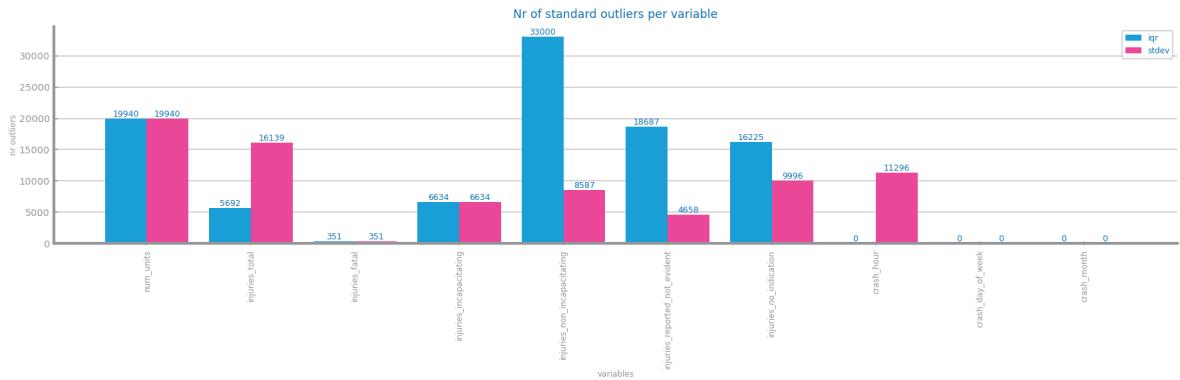


Figure 9: Outliers study dataset 1

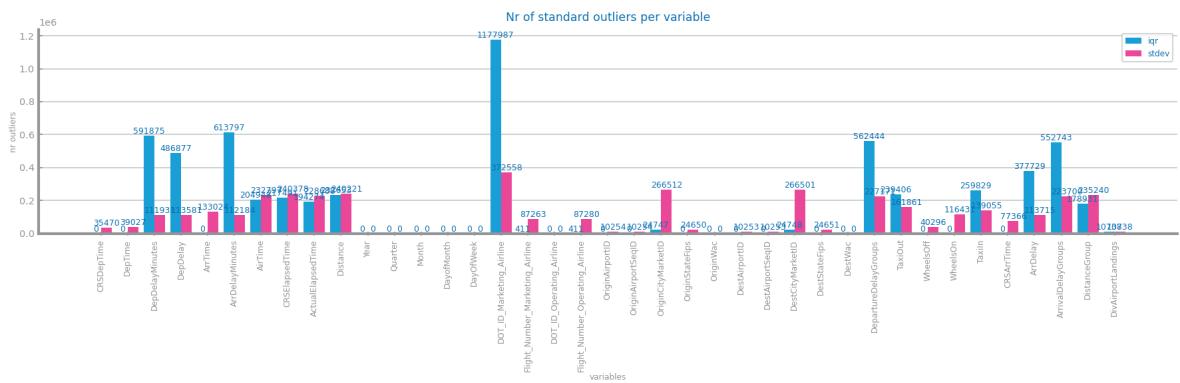


Figure 10: Outliers study dataset 2

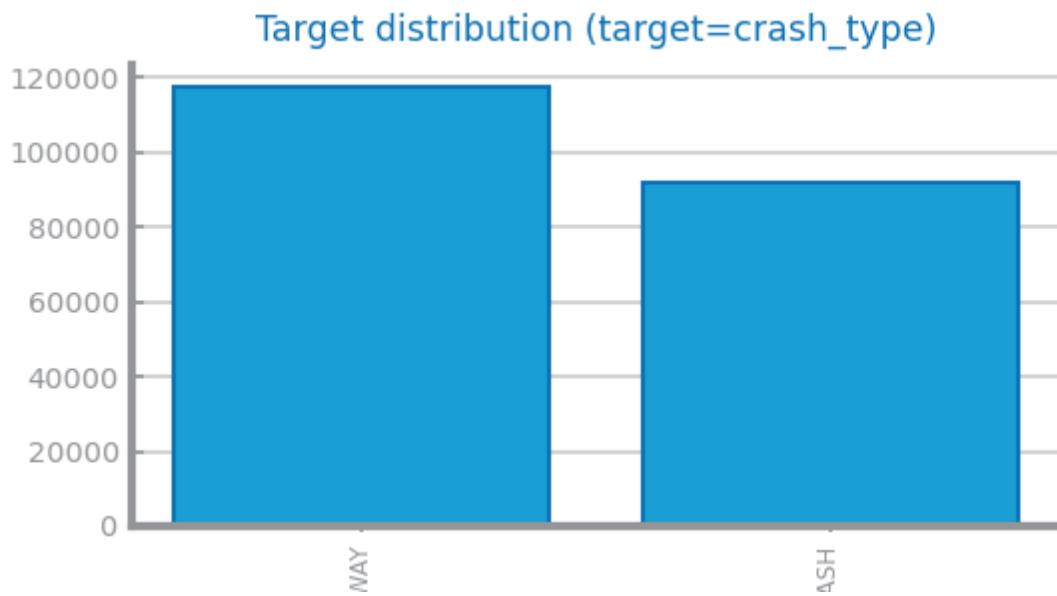


Figure 11: Class distribution for dataset 1

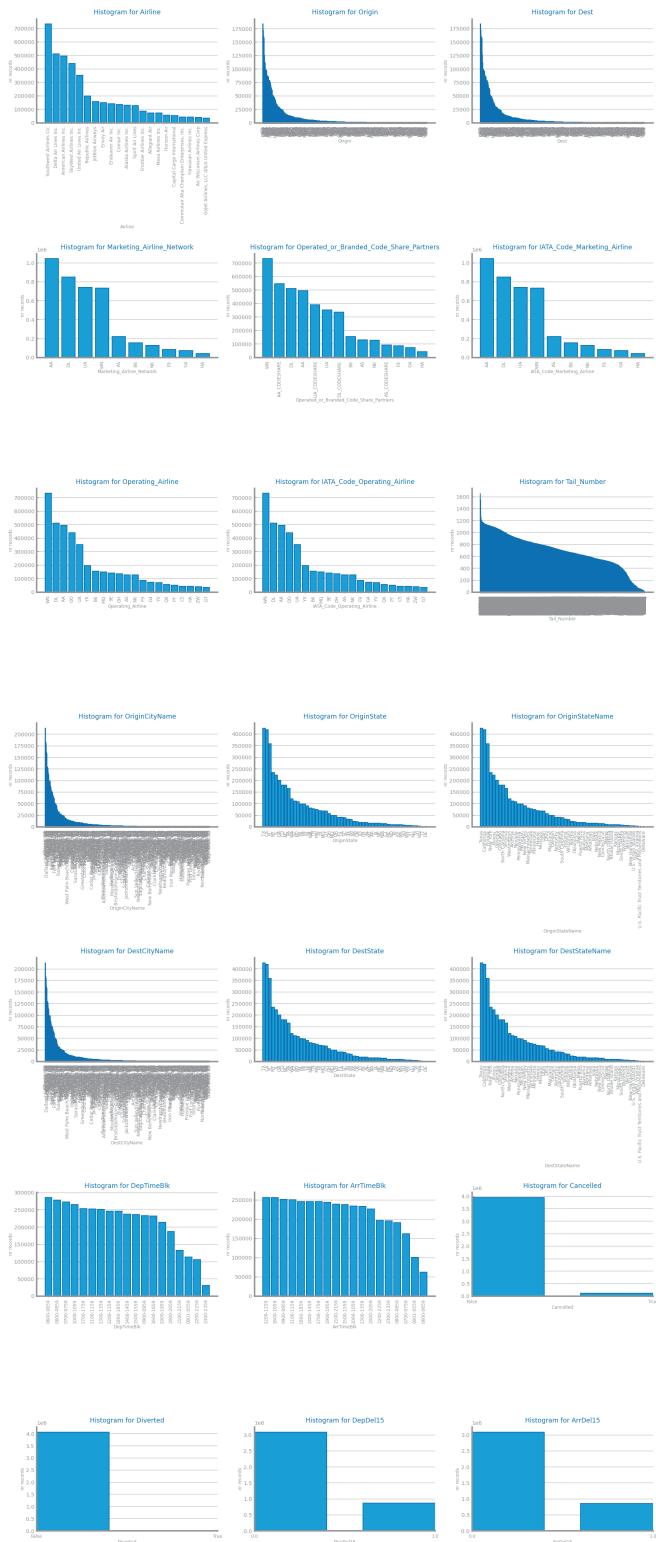


Figure 12: Class distribution for dataset 2

Data Granularity

The granularity analysis highlights significant variability across variables in both datasets. Categorical features such as weather, lighting conditions, airline, and airport present low to medium granularity, while temporal and numerical variables (hour, delays, distance) show higher granularity and dispersion. These differences suggest the need for appropriate grouping or discretization strategies to avoid sparsity and improve downstream analysis and modeling.

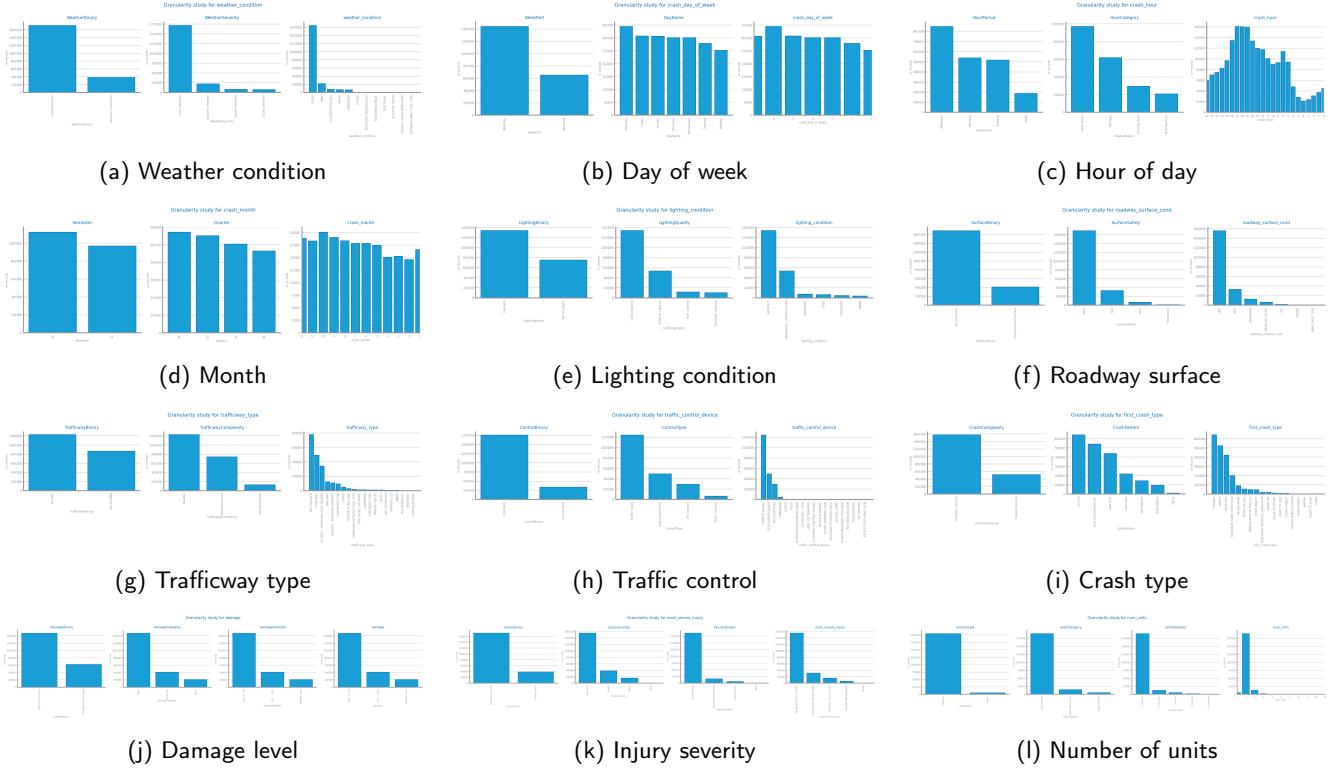


Figure 13: Granularity analysis for dataset 1

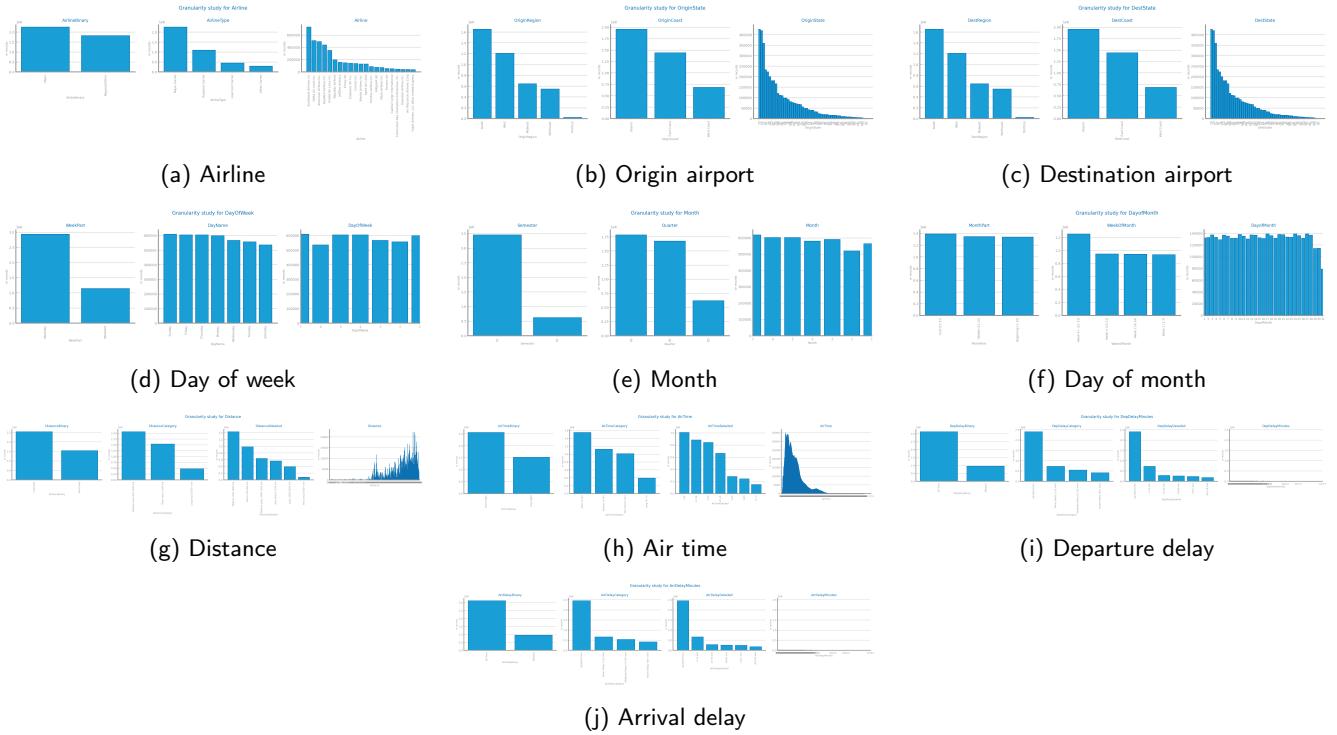


Figure 14: Granularity analysis for dataset 2

Data Sparsity

The sparsity analysis reveals uneven domain coverage across variables, with several categorical features dominated by a small number of frequent values and many rare ones. Correlation analysis shows limited strong dependencies between most variables, indicating sparse relationships in the feature space.



Figure 15: Sparsity analysis for dataset 1

Figure 16: Sparsity analysis for dataset 2 - [View on Google Drive](#)

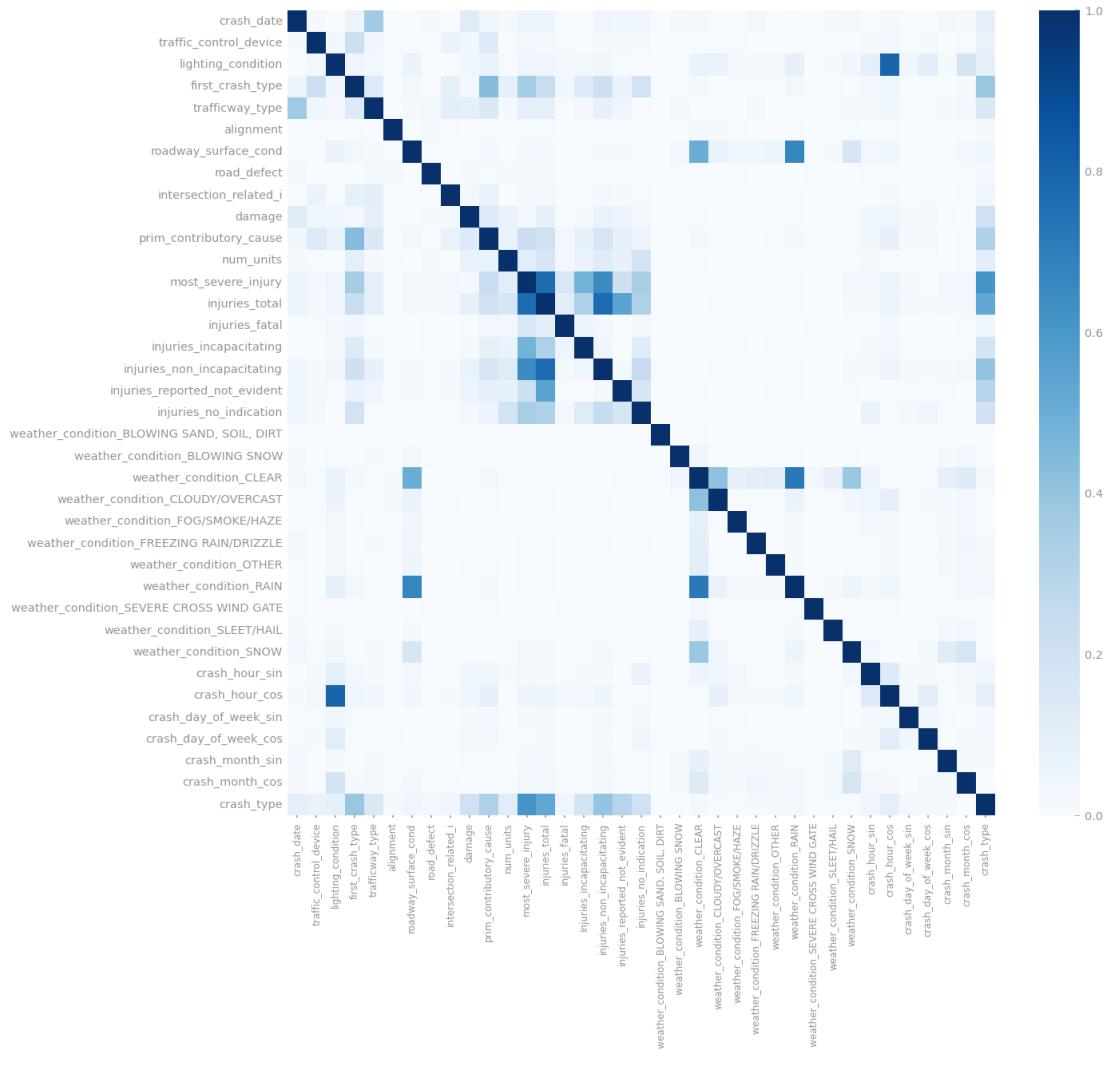


Figure 17: Correlation analysis for dataset 1

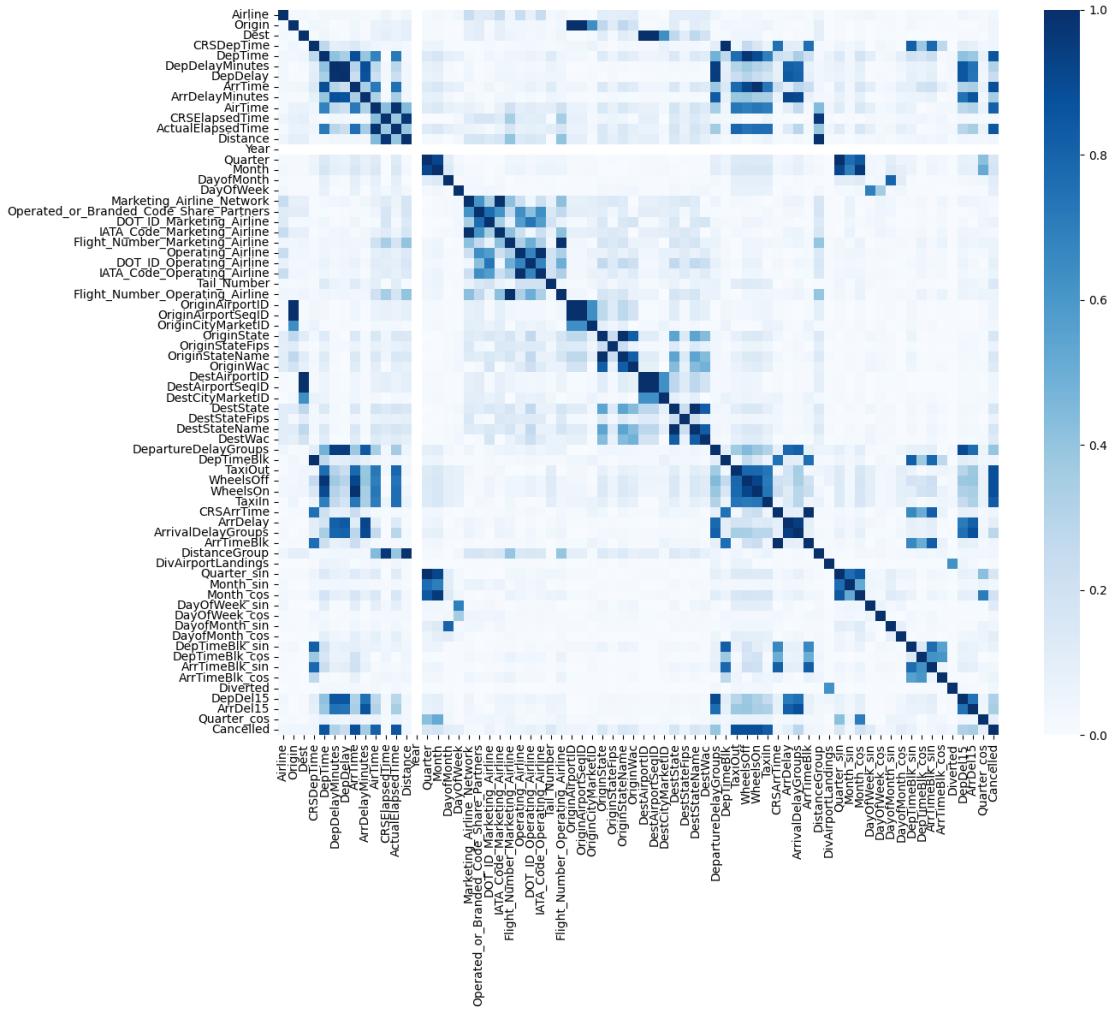


Figure 18: Correlation analysis for dataset 2

2 DATA PREPARATION

Variables Encoding

Shall contain all relevant information respecting to the transformation of variables. The list of variables under each one of the transformations, shall be presented. If not applied explain the reason for that, based on data characteristics. **Shall not exceed 500 characters for each dataset.**

Missing Value Imputation

Shall contain all relevant information and charts respecting to missing values imputation, such as the choices made and the impact of the different approaches on modelling results. Shall also clearly reveal the approach selected to proceed with the processing. If not applied explain the reason for that, based on data characteristics. **Shall not exceed 500 characters.**

Figure 19: Missing values imputation results with different approaches for dataset 1

Figure 20: Missing values imputation results with different approaches for dataset 2

Outliers Treatment

Shall contain all relevant information and charts respecting to outliers imputation, such as the choices made and the impact of the different approaches on modelling results. Shall also clearly reveal the approach selected to proceed with the processing. If not applied explain the reason for that, based on data characteristics. **Shall not exceed 500 characters.**

Figure 21: Outliers imputation results with different approaches for dataset 1

Figure 22: Outliers imputation results with different approaches for dataset 2

Scaling

Shall contain all relevant information and charts respecting to scaling transformation, such as the choices made and the impact of the different approaches on modelling results. Shall also clearly reveal the approach selected to proceed with the processing. If not applied explain the reason for that, based on data characteristics. **Shall not exceed 200 characters.**

Figure 23: Scaling results with different approaches for dataset 1

Figure 24: Scaling results with different approaches for dataset 2

Balancing

Shall contain all relevant information and charts respecting to balancing transformation, such as the choices made and the impact of the different approaches on modelling results. Shall also clearly reveal the approach selected to proceed with the processing. If not applied explain the reason for that, based on data characteristics. **Shall not exceed 500 characters.**

Figure 25: Balancing results with different approaches for dataset 1

Figure 26: Balancing results with different approaches for dataset 2

Feature Selection

Shall contain all relevant information and charts respecting to feature selection based on filtering out redundant (based on correlation) and relevant (based on variation) variables. The different choices and their impact on the modelling results shall be presented and explained. Should also clearly reveal the approach selected to proceed with the processing. All explanations shall be based on data characteristics. **Shall not exceed 500 characters.**

Figure 27: Feature selection of redundant variables results with different parameters for dataset 1

Figure 28: Feature selection of redundant variables results with different parameters for dataset 2

Figure 29: Feature selection of relevant variables results with different parameters for dataset 1 (variance study)

Figure 30: Feature selection of relevant variables results with different parameters for dataset 2 (variance study)

Feature Extraction (optional)

Shall contain all relevant information and charts respecting to feature extraction, in particular PCA. The different choices and their impact on the modelling results shall be presented and explained. **Shall not exceed 200 characters.**

Figure 31: Principal components analysis and feature extraction results for dataset 1

Figure 32: Principal components analysis and feature extraction results for dataset 2

Additional Feature Generation (if done)

Shall contain all relevant information and charts respecting to feature generation. The different choices and their impact on the modelling results shall be presented and explained. Shall summarise all variables generated and the formula used to derive them (in a table). **Shall not exceed 200 characters.**

Figure 33: Feature generation results for dataset 1

Figure 34: Feature generation results for dataset 2

3 MODELS' EVALUATION

Shall be used to point out any important decision taken during the training, including training strategy and evaluation measures used. **Shall not exceed 500 characters.**

Naïve Bayes

Shall be used to present the results achieved with each one of Naïve Bayes implementations, comparing and proposing explanations for them. If any of the implementations is not used, a justification for it shall be presented. Shall be used to present the evaluation of the best model achieved. **Shall not exceed 300 characters.**

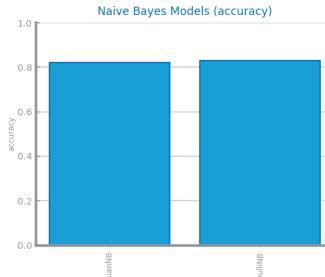


Figure 35: Naïve Bayes alternatives comparison for dataset 1

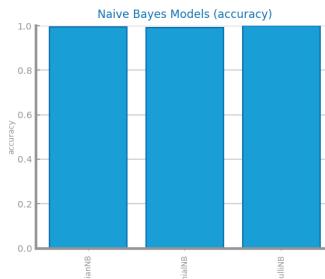


Figure 36: Naïve Bayes alternative comparison for dataset 2

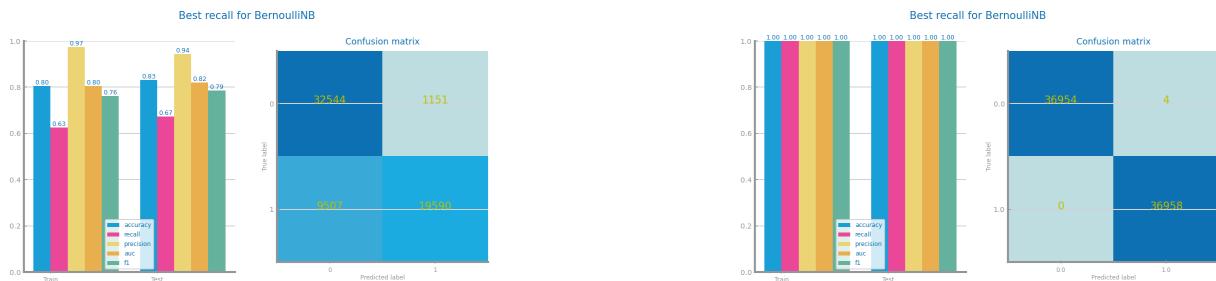


Figure 37: Naïve Bayes best model results for dataset 1 (left) and dataset 2 (right)

KNN

Shall be used to present the results achieved through different similarity measures and KNN parameterisations. The results shall be compared and explanations for them shall be presented. The justification for the chosen similarity measures shall be presented. Shall be used to address the overfitting phenomenon, studying the conditions under which models face it. Shall be used to present the evaluation of the best model achieved. **Shall not exceed 500 characters.**

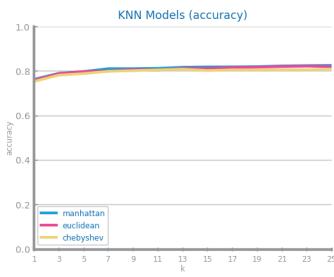


Figure 38: KNN different parameterisations comparison for dataset 1

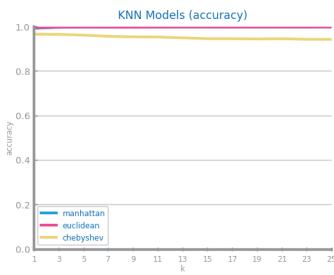


Figure 39: KNN different parameterisations comparison for dataset 2

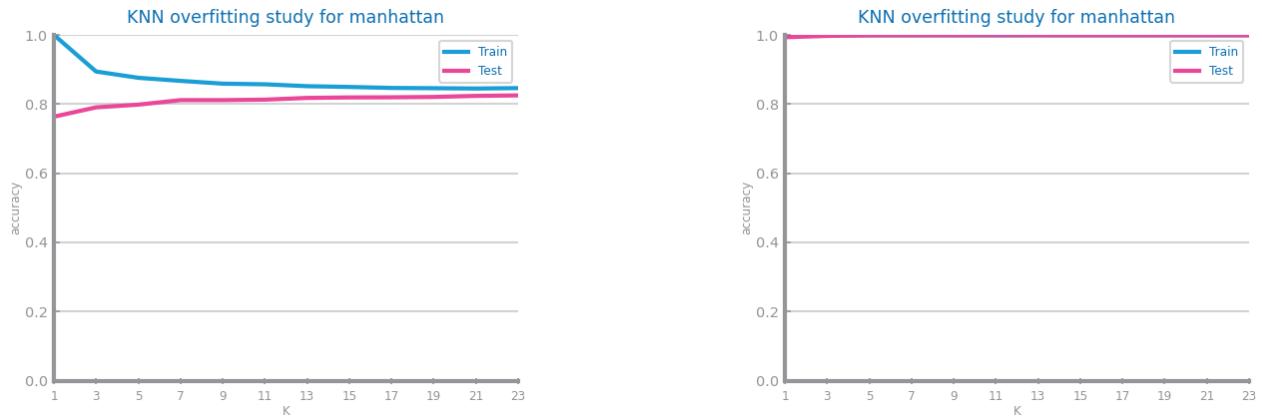


Figure 40: KNN overfitting analysis for dataset 1 (left) and dataset 2 (right)

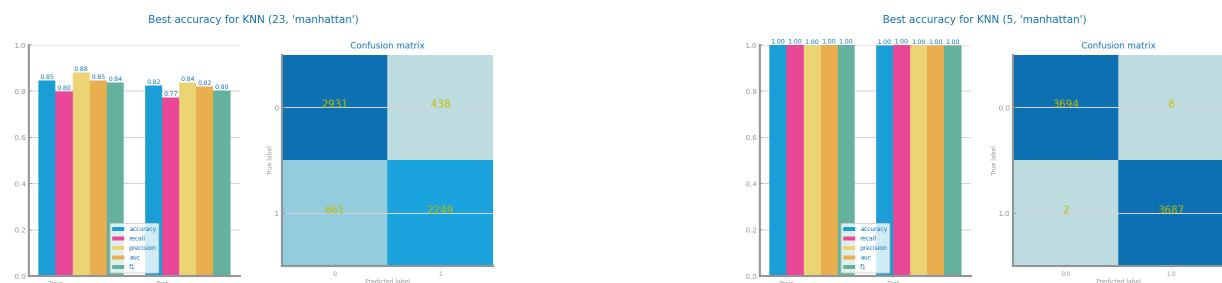


Figure 41: KNN best model results for dataset 1 (left) and dataset 2 (right)

Logistic Regression

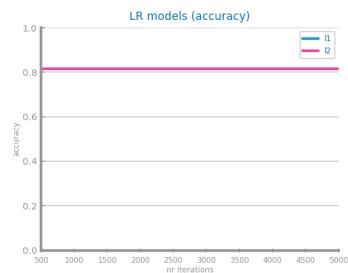


Figure 42: Logistic Regression different parameterisations comparison for dataset 1

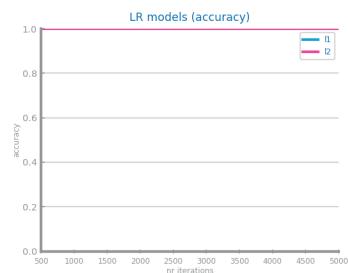


Figure 43: Logistic Regression different parameterisations comparison for dataset 2

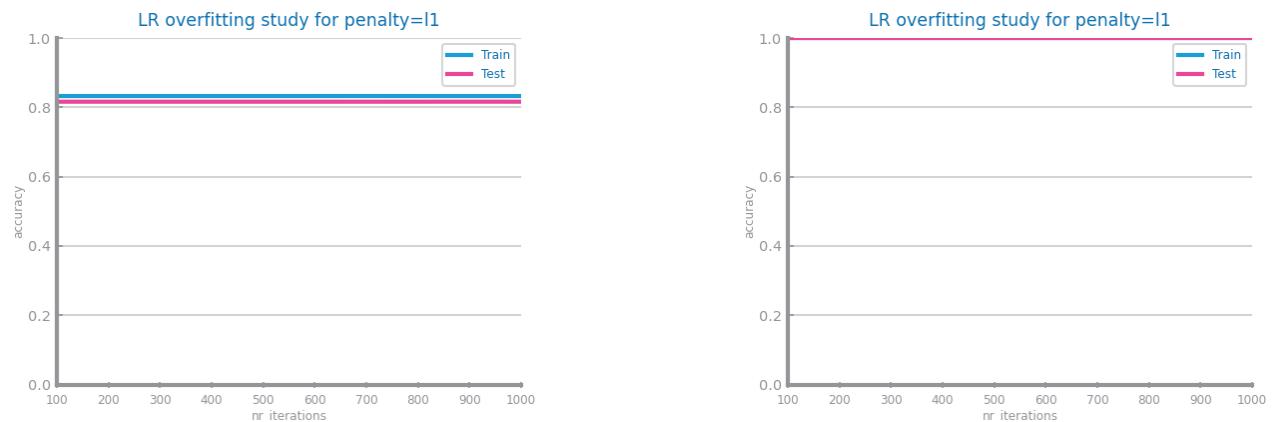


Figure 44: Logistic Regression overfitting analysis for dataset 1 (left) and dataset 2 (right)

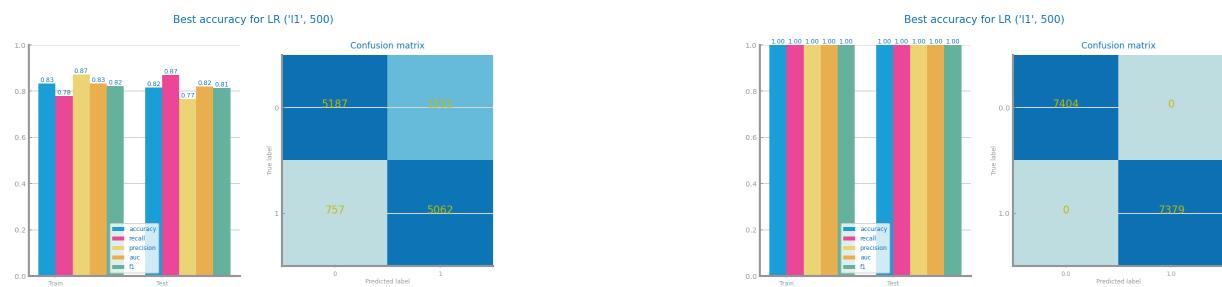


Figure 45: Logistic Regression best model results for dataset 1 (left) and dataset 2 (right)

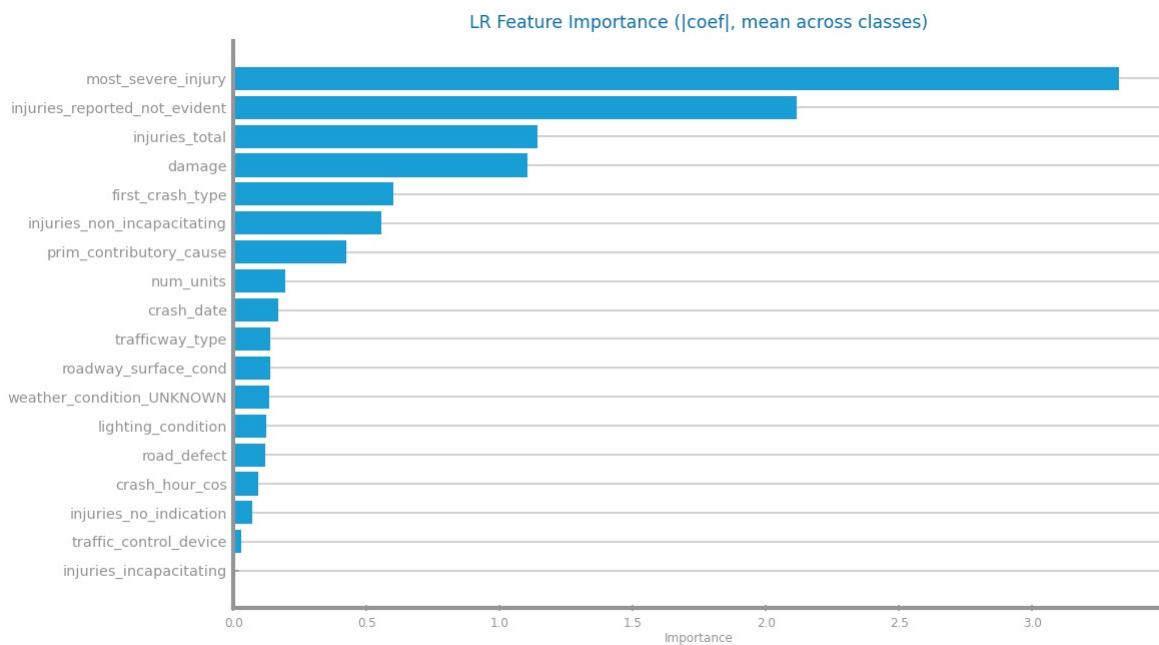


Figure 46: Logistic Regression feature importance for dataset 1

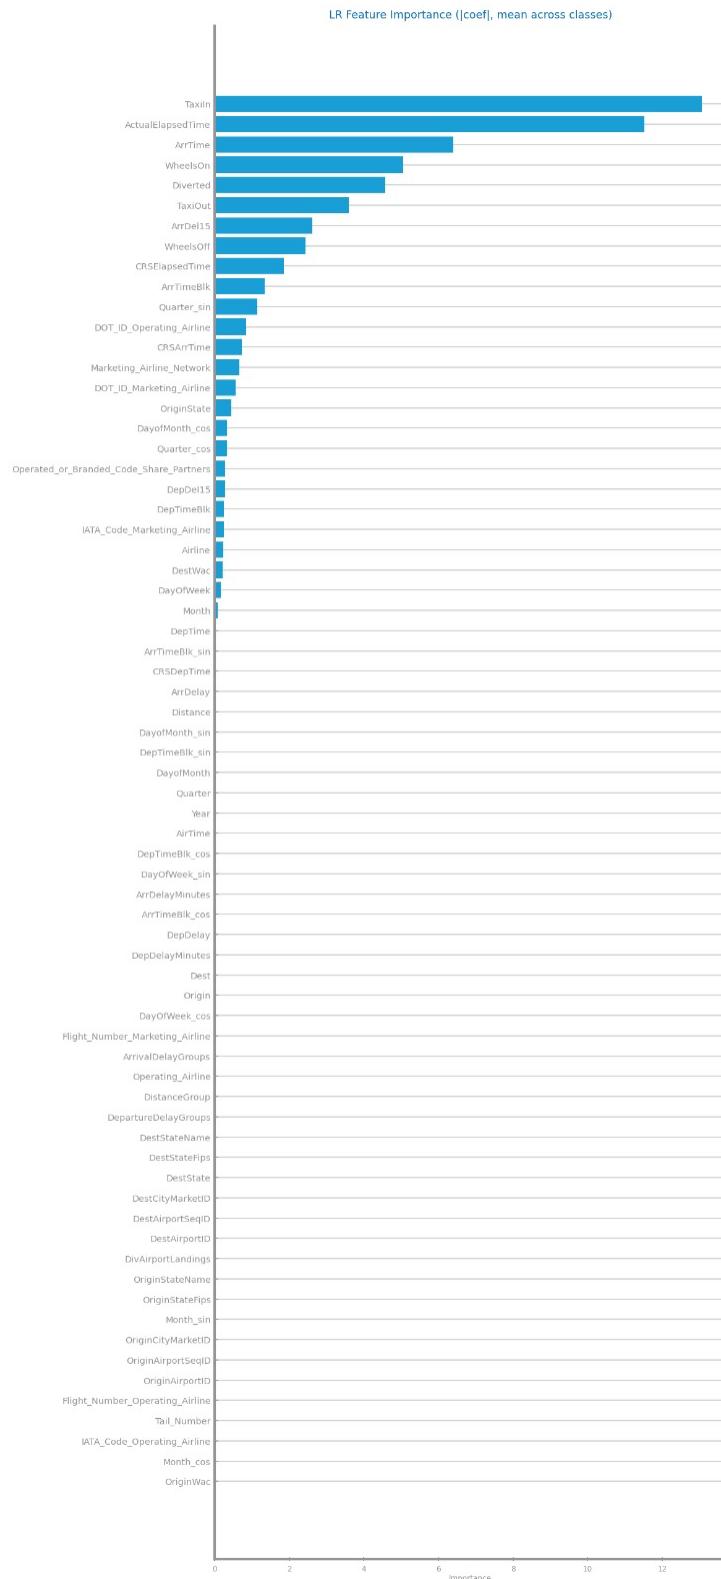


Figure 47: Logistic Regression feature importance for dataset 2

Decision Trees

Shall be used to present the results achieved through different parameterisations for the train of decision trees. The results shall be compared and explanations for them shall be presented. Shall be used to address the overfitting phenomenon, studying the conditions under which models face it. Shall be used to present the evaluation of the best model achieved. Shall be used to present the best tree achieved and its succinct description. **Shall not exceed 500 characters.**

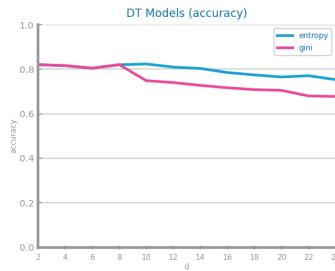


Figure 48: Decision Trees different parameterisations comparison for dataset 1

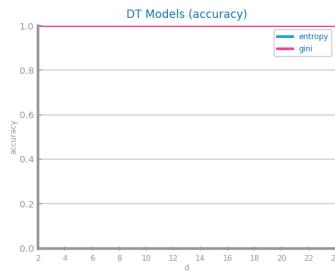


Figure 49: Decision Trees different parameterisations comparison for dataset 2

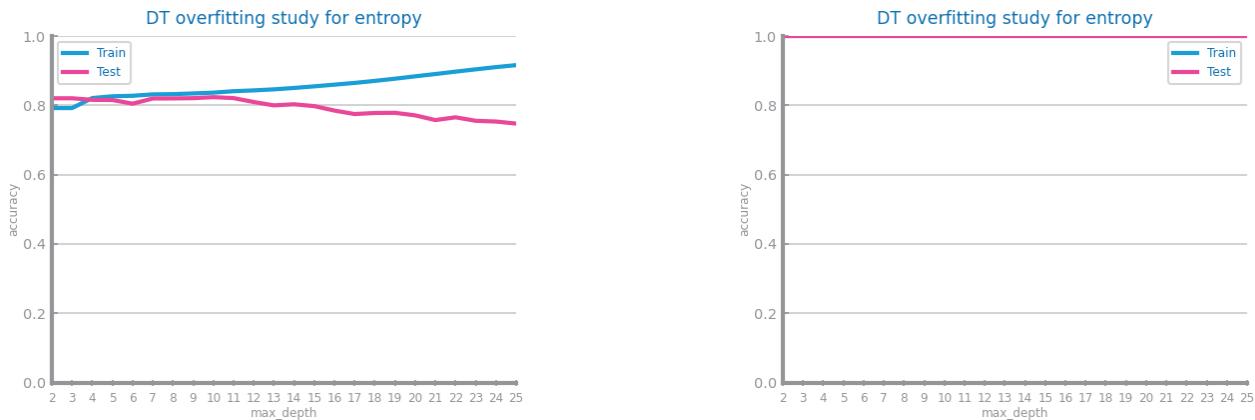


Figure 50: Decision Trees overfitting analysis for dataset 1 (left) and dataset 2 (right)

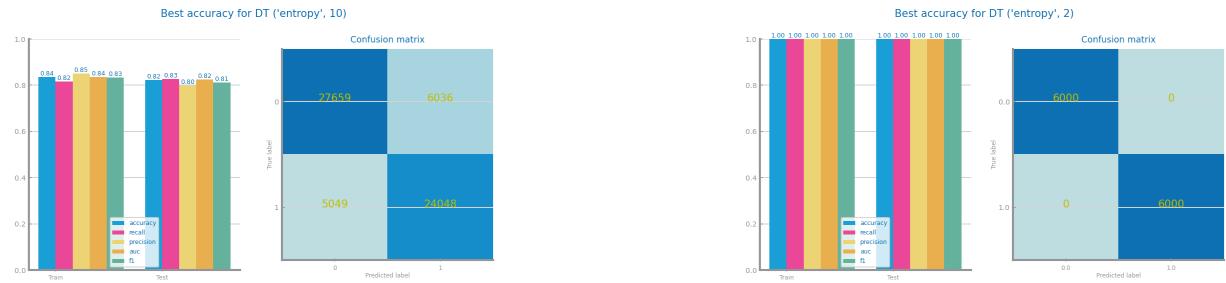


Figure 51: Decision trees best model results for dataset 1 (left) and dataset 2 (right)

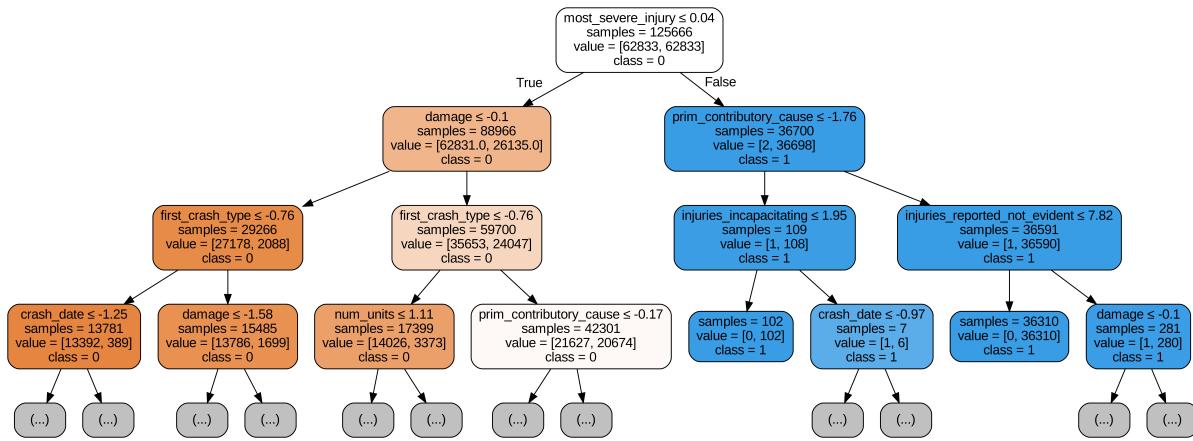


Figure 52: Best tree for dataset 1

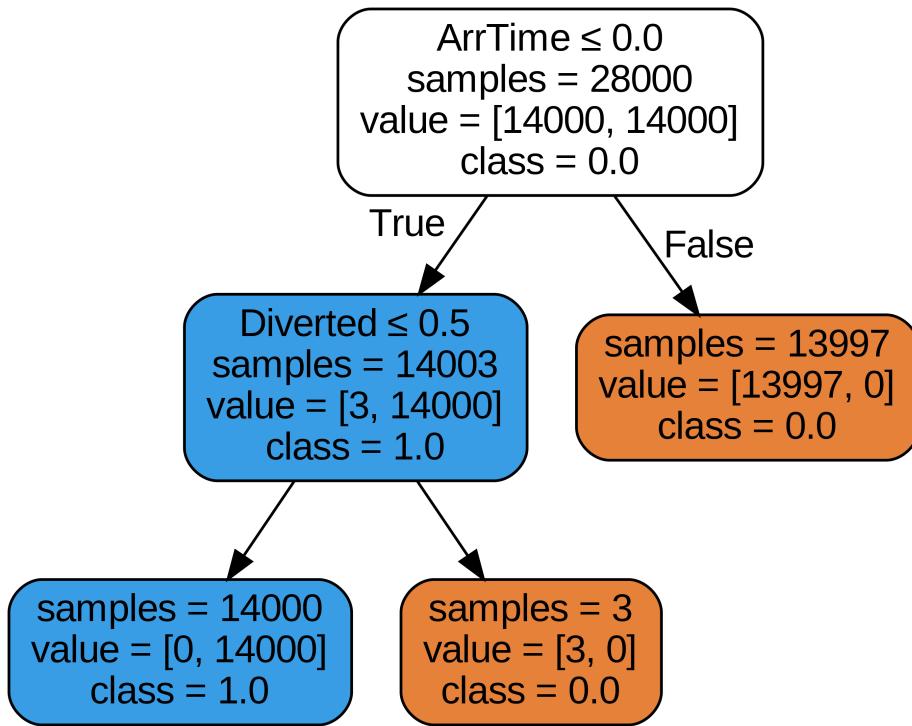


Figure 53: Best tree for dataset 2

Random Forests

Shall be used to present the results achieved through different parameterisations for the train of random forests. The results shall be compared and explanations for them shall be presented. Shall be used to address the overfitting phenomenon, studying the conditions under which models face it. Shall be used to present the evaluation of the best model achieved. May be used to present the most important variables in the model. **Shall not exceed 500 characters.**

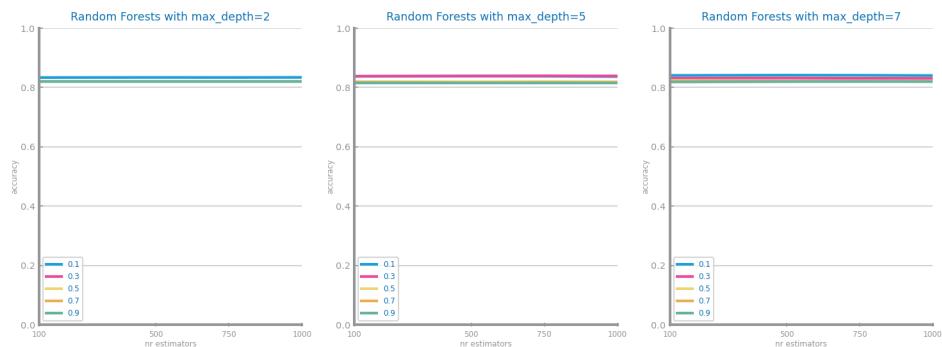


Figure 54: Random Forests different parameterisations comparison for dataset 1

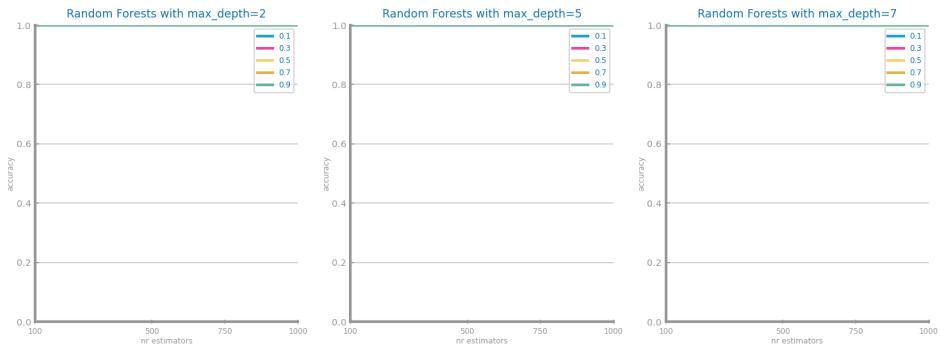


Figure 55: Random Forests different parameterisations comparison for dataset 2

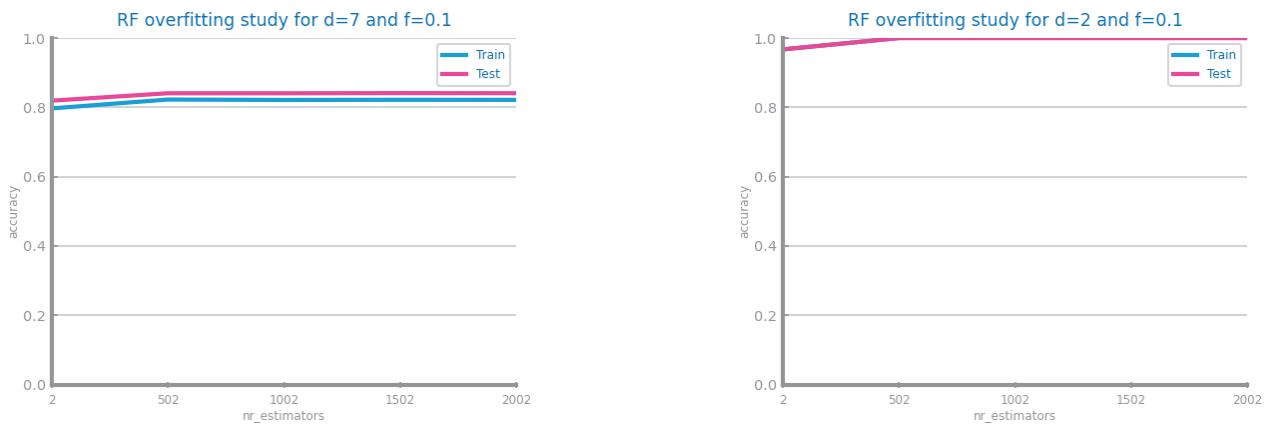


Figure 56: Random Forests overfitting analysis for dataset 1 (left) and dataset 2 (right)

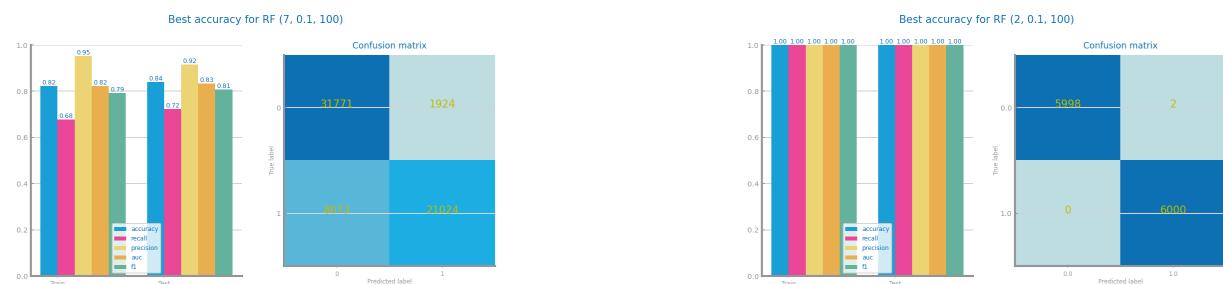


Figure 57: Random Forests best model results for dataset 1 (left) and dataset 2 (right)

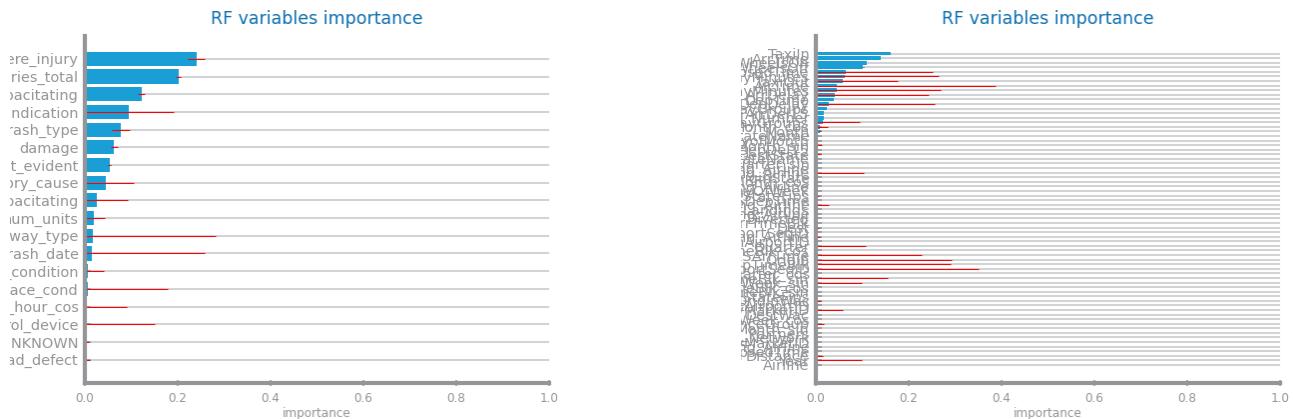


Figure 58: Random Forests variables importance for dataset 1 (left) and dataset 2 (right)

Gradient Boosting

Shall be used to present the results achieved through different parameterisations for the train of gradient boosting. The results shall be compared and explanations for them shall be presented. Shall be used to address the overfitting phenomenon, studying the conditions under which models face it. Shall be used to present the evaluation of the best model achieved. May be used to present the most important variables in the model. **Shall not exceed 500 characters.**

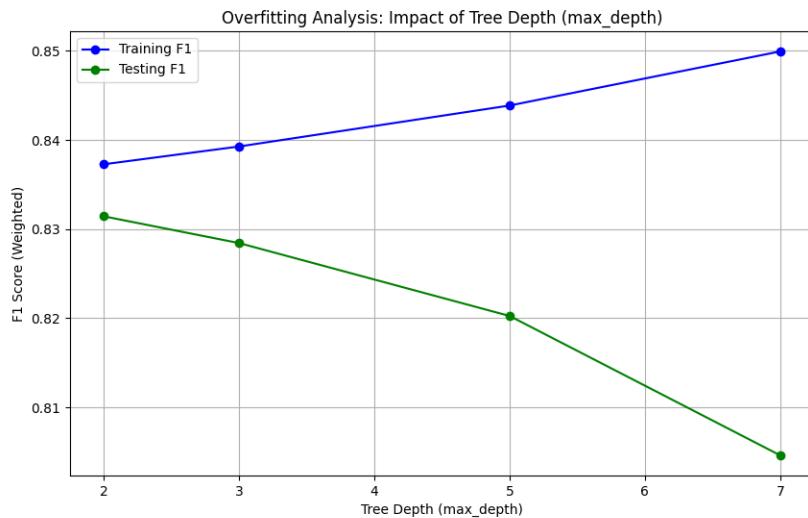


Figure 59: Gradient boosting different parameterisations comparison for dataset 1

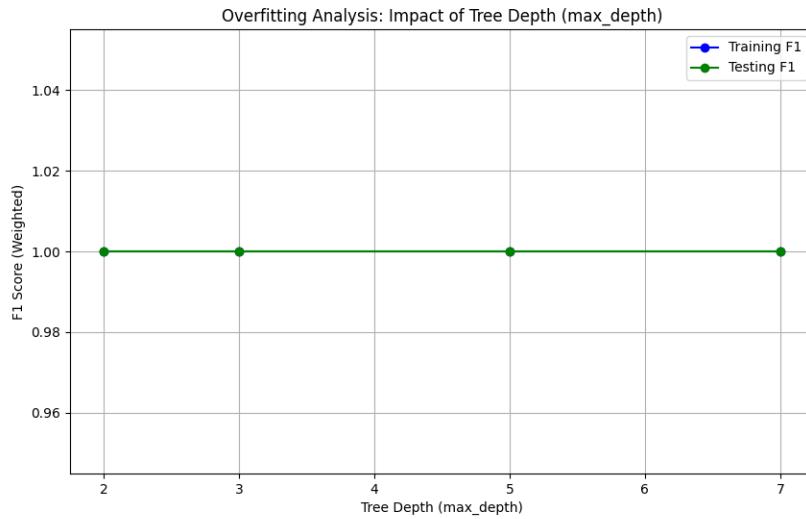


Figure 60: Gradient boosting different parameterisations comparison for dataset 2

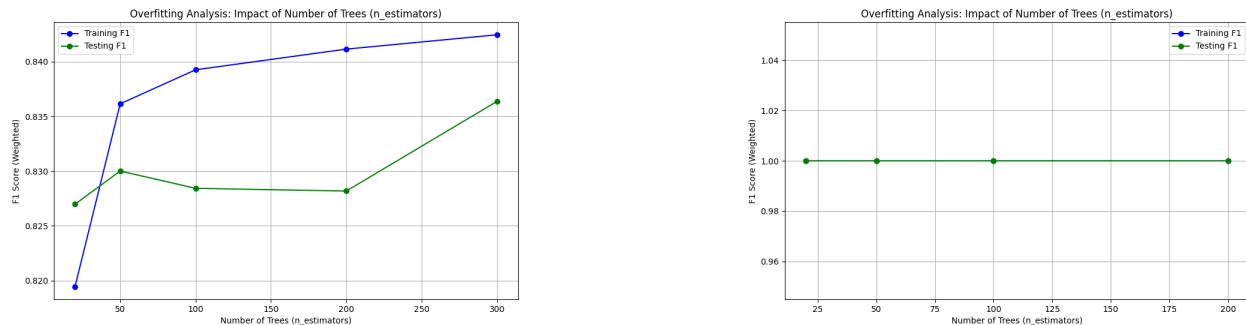


Figure 61: Gradient boosting overfitting analysis for dataset 1 (left) and dataset 2 (right)

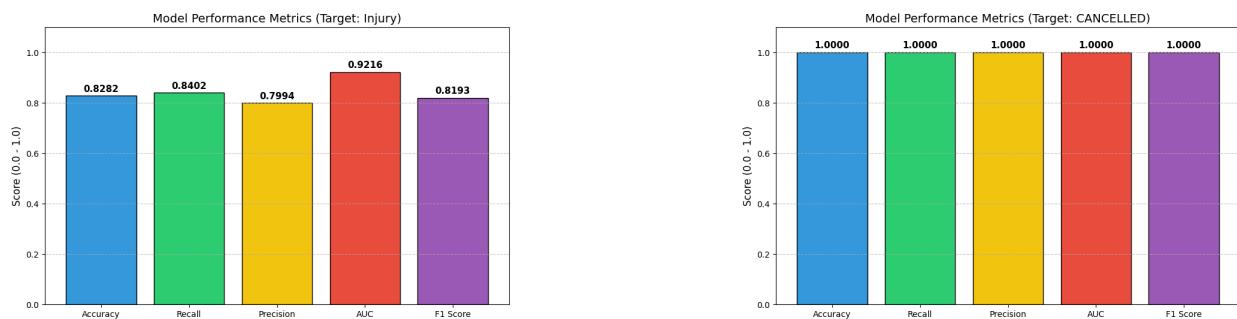


Figure 62: Gradient boosting best model results for dataset 1 (left) and dataset 2 (right)

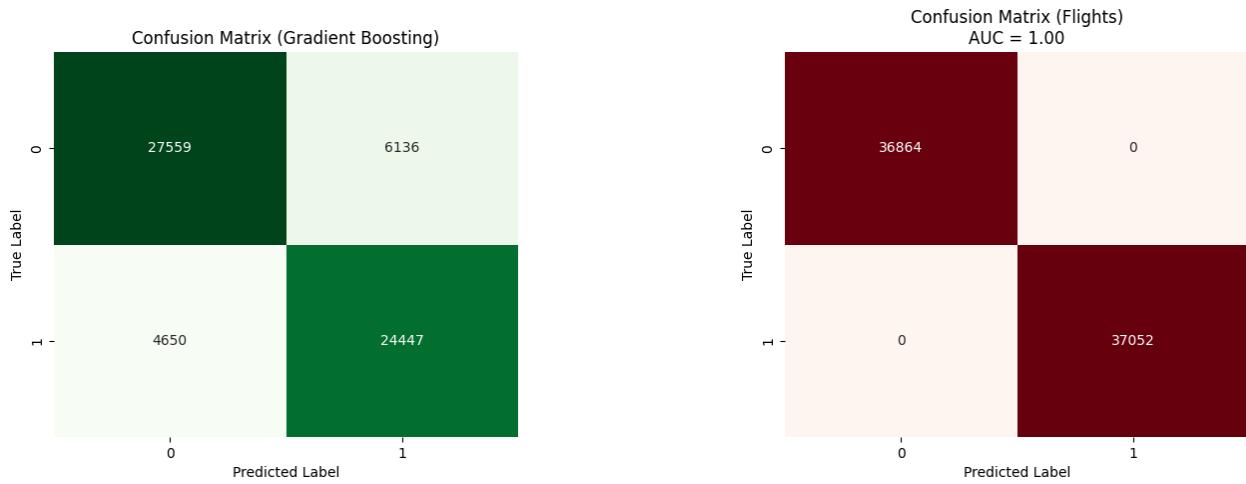


Figure 63: Gradient boosting variables importance for dataset 1 (left) and dataset 2 (right)

Multi-Layer Perceptrons

Shall be used to present the results achieved through different parameterisations for the train of MLPs. The results shall be compared and explanations for them shall be presented. Shall be used to address the overfitting phenomenon, studying the conditions under which models face it. Shall be used to present the evaluation of the best model achieved. **Shall not exceed 500 characters.**

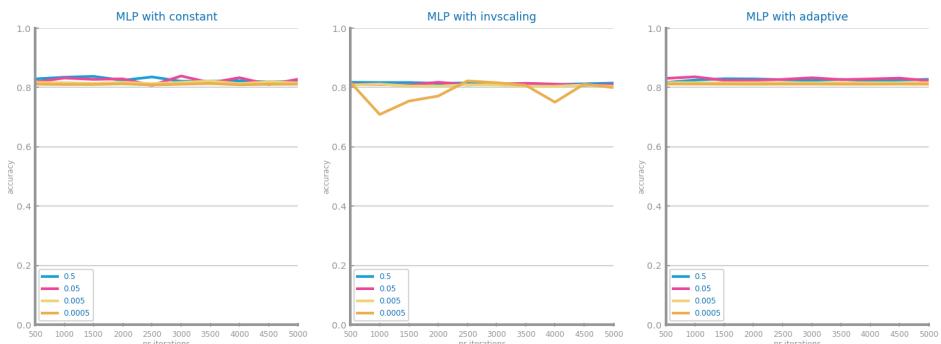


Figure 64: MLP different parameterisations comparison for dataset 1

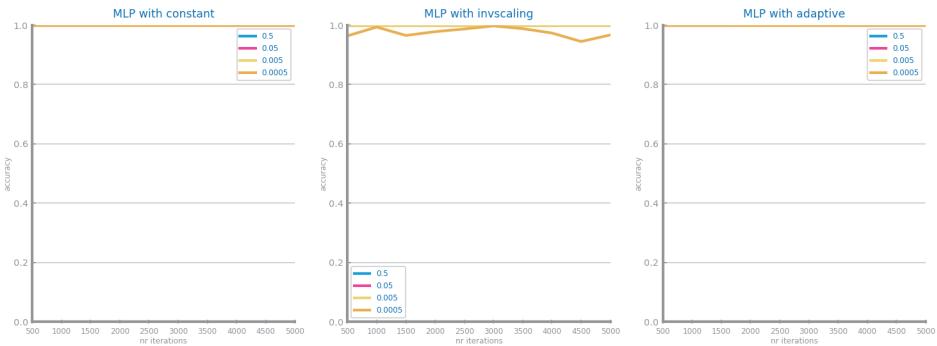


Figure 65: MLP different parameterisations comparison for dataset 2

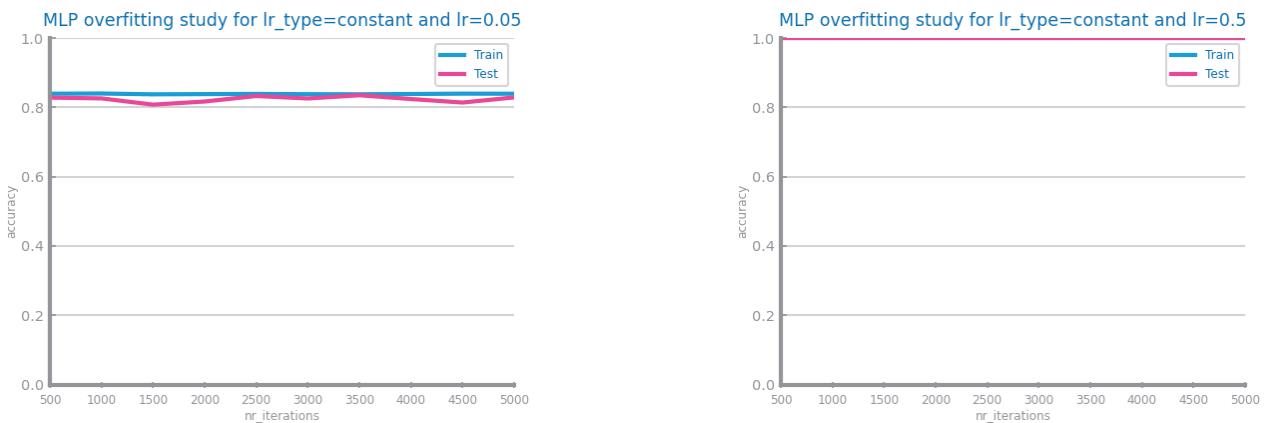


Figure 66: MLP overfitting analysis for dataset 1 (left) and dataset 2 (right)

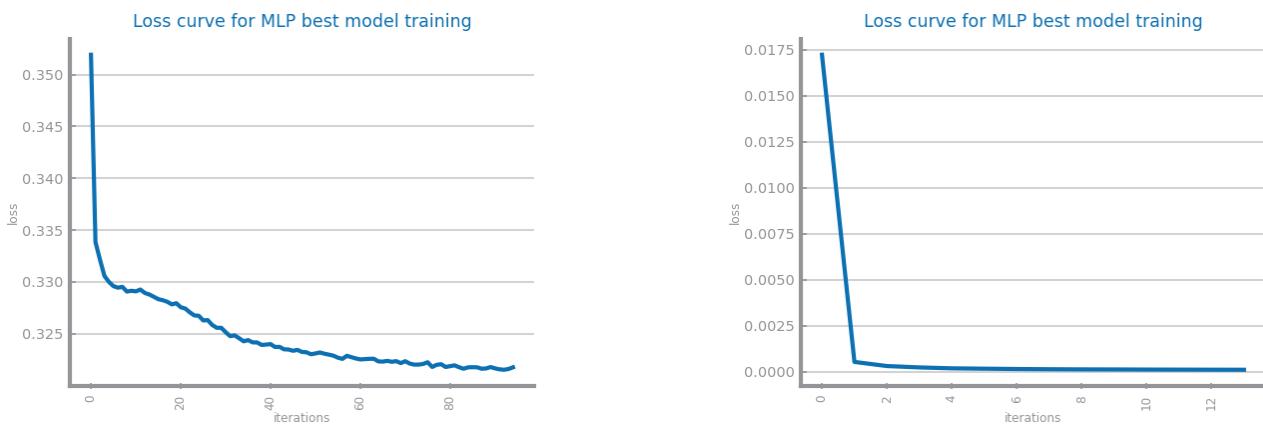


Figure 67: Loss curve analysis for dataset 1 (left) and dataset 2 (right)

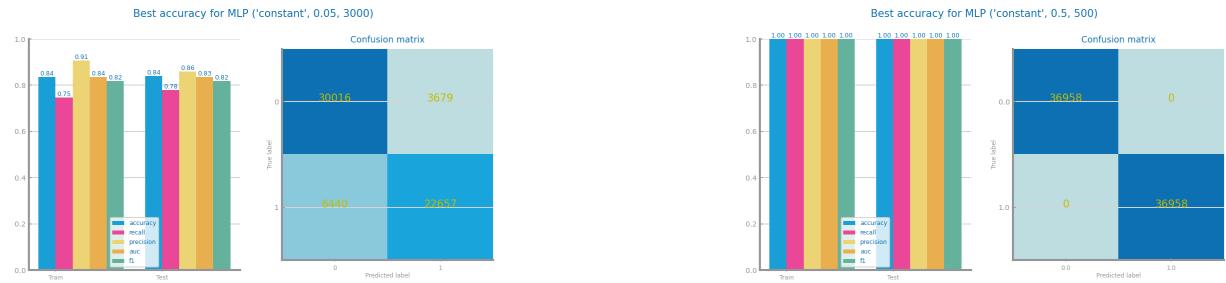


Figure 68: MLP best model results for dataset 1 (left) and dataset 2 (right)

4 CRITICAL ANALYSIS

Shall be used to present a summary of the results achieved with the different modelling techniques, and the impact of the different preparation tasks on their performance. A cross-analysis of the different models may also be presented, identifying the most relevant variables common to all of them (when possible) and the relation among the patterns identified within the different classifiers. A critical assessment of the best models shall be presented, clearly stating if the models seem to be good enough for the problem at hand. **Additional charts may be presented here. Shall not exceed 2000 characters.**

TIME SERIES ANALYSIS

5 DATA PROFILING

Data Dimensionality and Granularity

The time series was analyzed at three temporal granularities: 15-minute intervals (most atomic), daily aggregation, and weekly aggregation. The most granular level captures short-term variability and peaks, while daily and weekly granularities smooth fluctuations and highlight medium- and long-term trends.

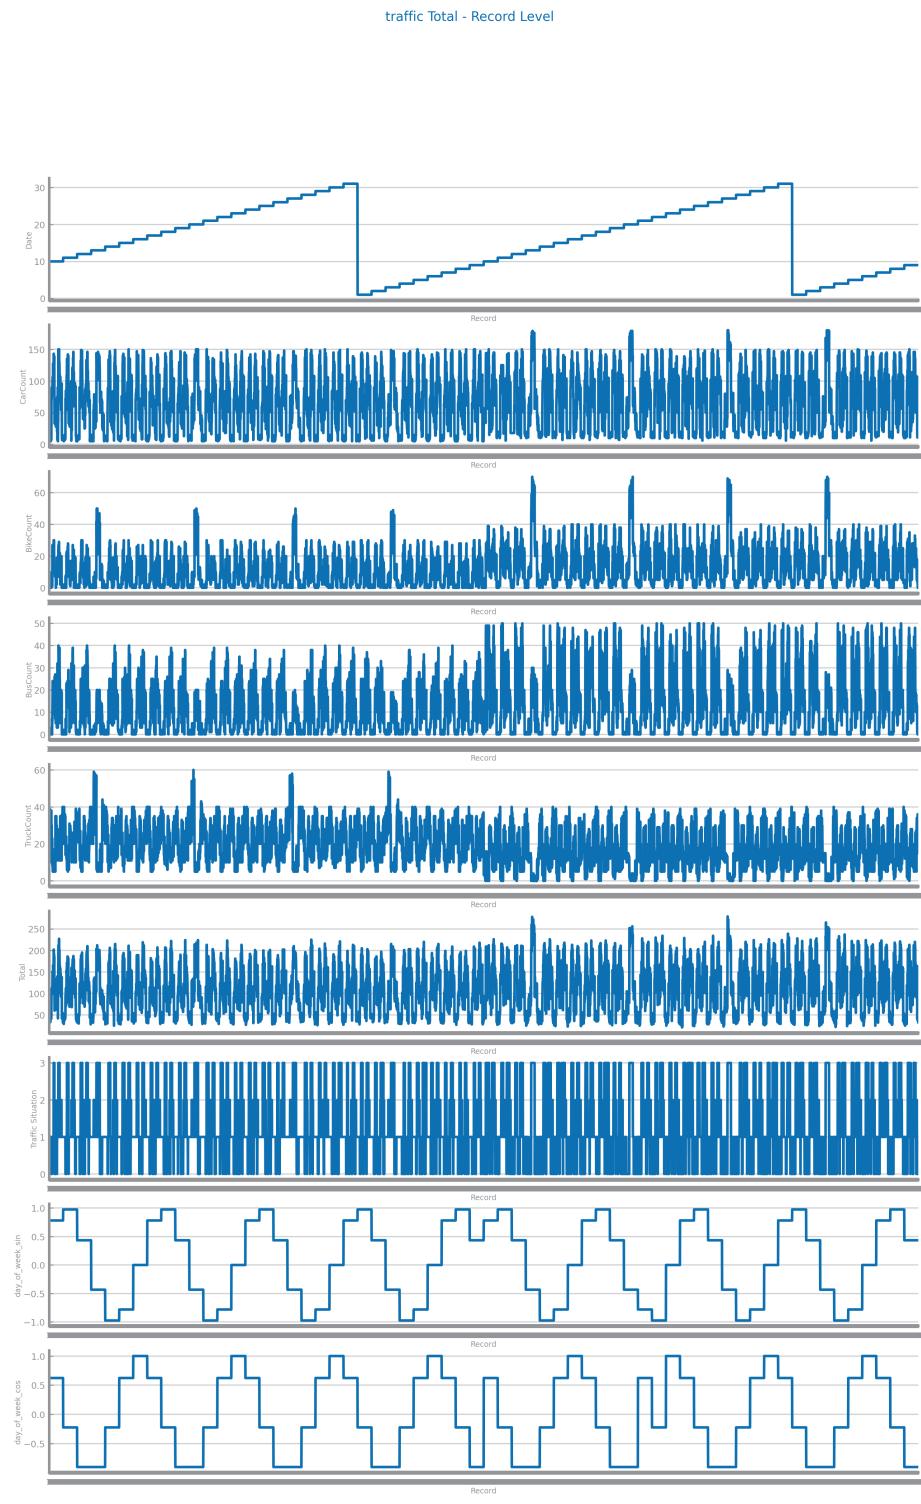


Figure 69: Time series 1 at the most granular detail

traffic Total - Daily

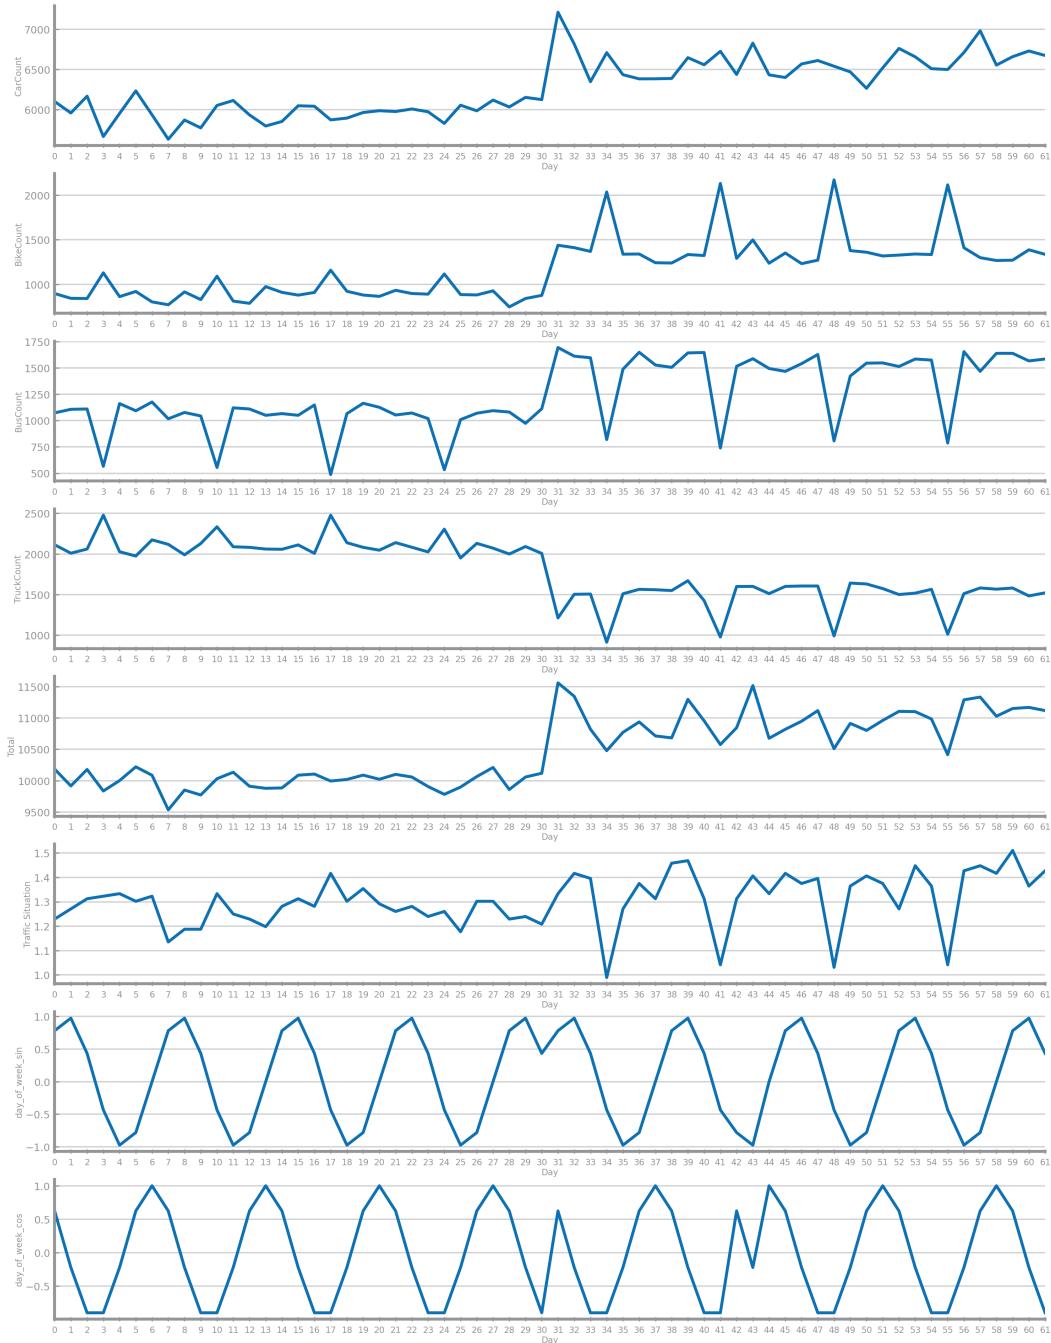


Figure 70: Time series 1 at the second chosen granularity

traffic Total - Weekly

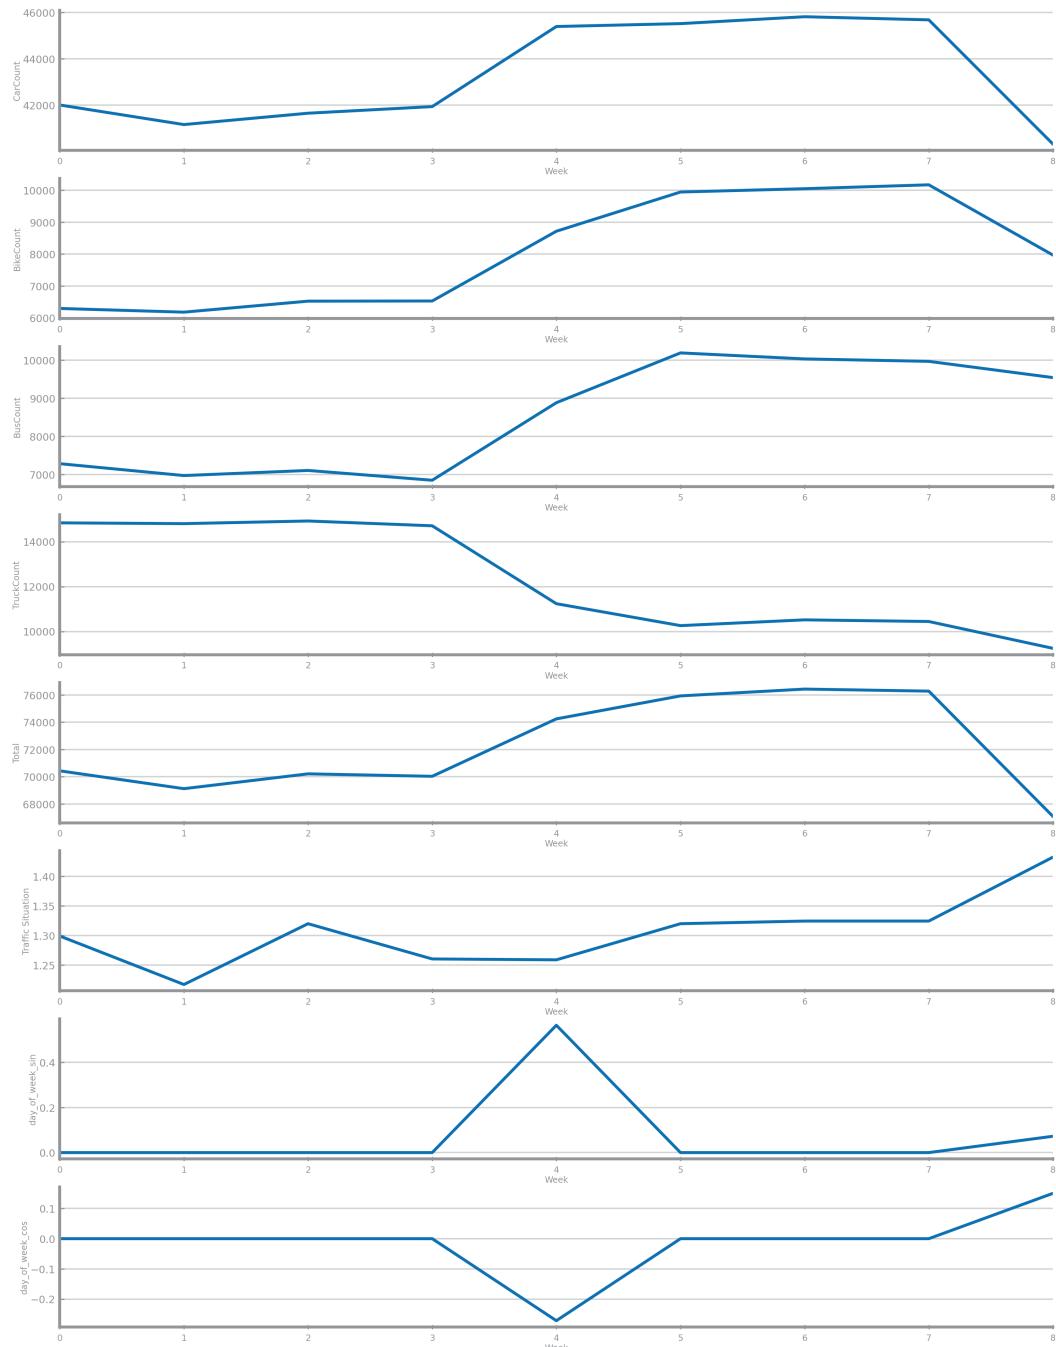


Figure 71: Time series 1 at the third chosen granularity

Data Distribution

The distribution analysis across granularities shows that aggregation reduces variability and outliers. The 15-minute series presents higher dispersion and skewness, while hourly and daily series exhibit smoother distributions.

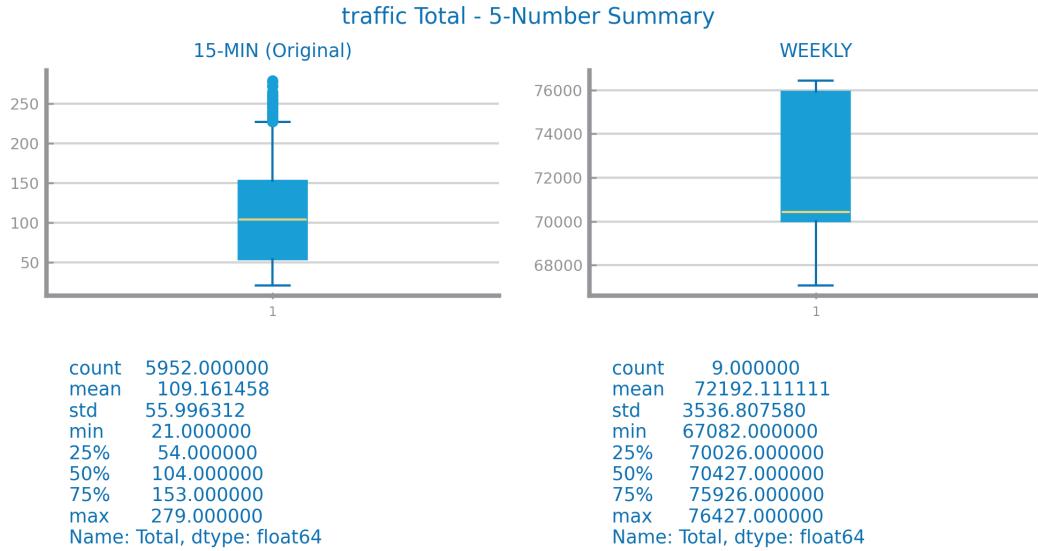


Figure 72: Boxplot(s) for time series 1

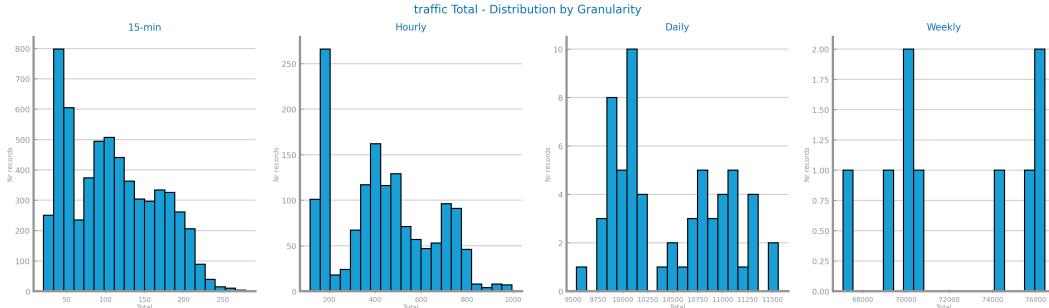


Figure 73: Histogram(s) for time series 1



Figure 74: Autocorrelation lag-plots for original time series 1

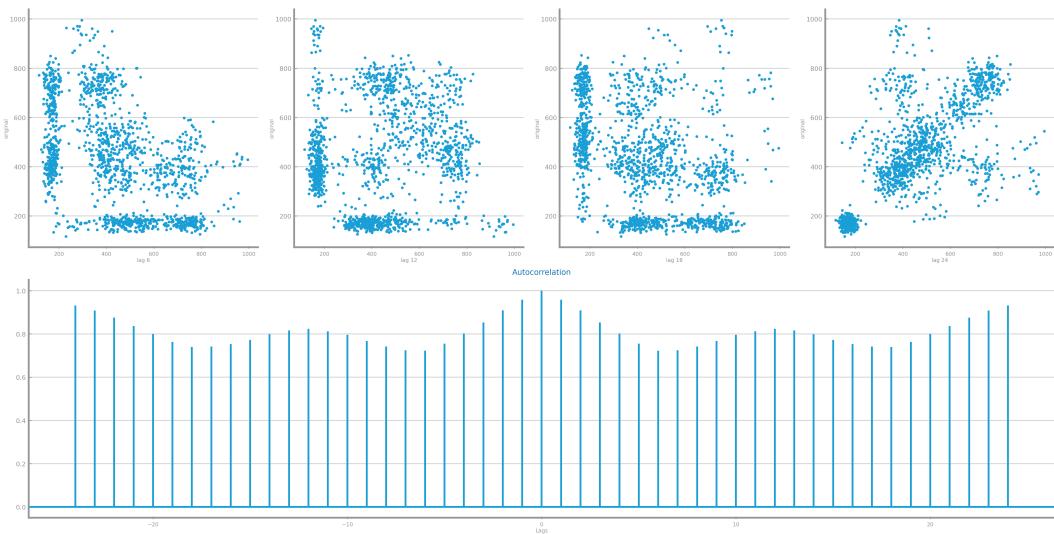


Figure 75: Autocorrelation correlogram for original time series 1

Data Stationarity

Stationarity analysis shows that the series is stationary at finer granularities but loses stationarity when aggregated daily. Augmented Dickey-Fuller Test confirm that trend components become more pronounced at coarser granularities, indicating the need for detrending or differencing when modeling aggregated series.

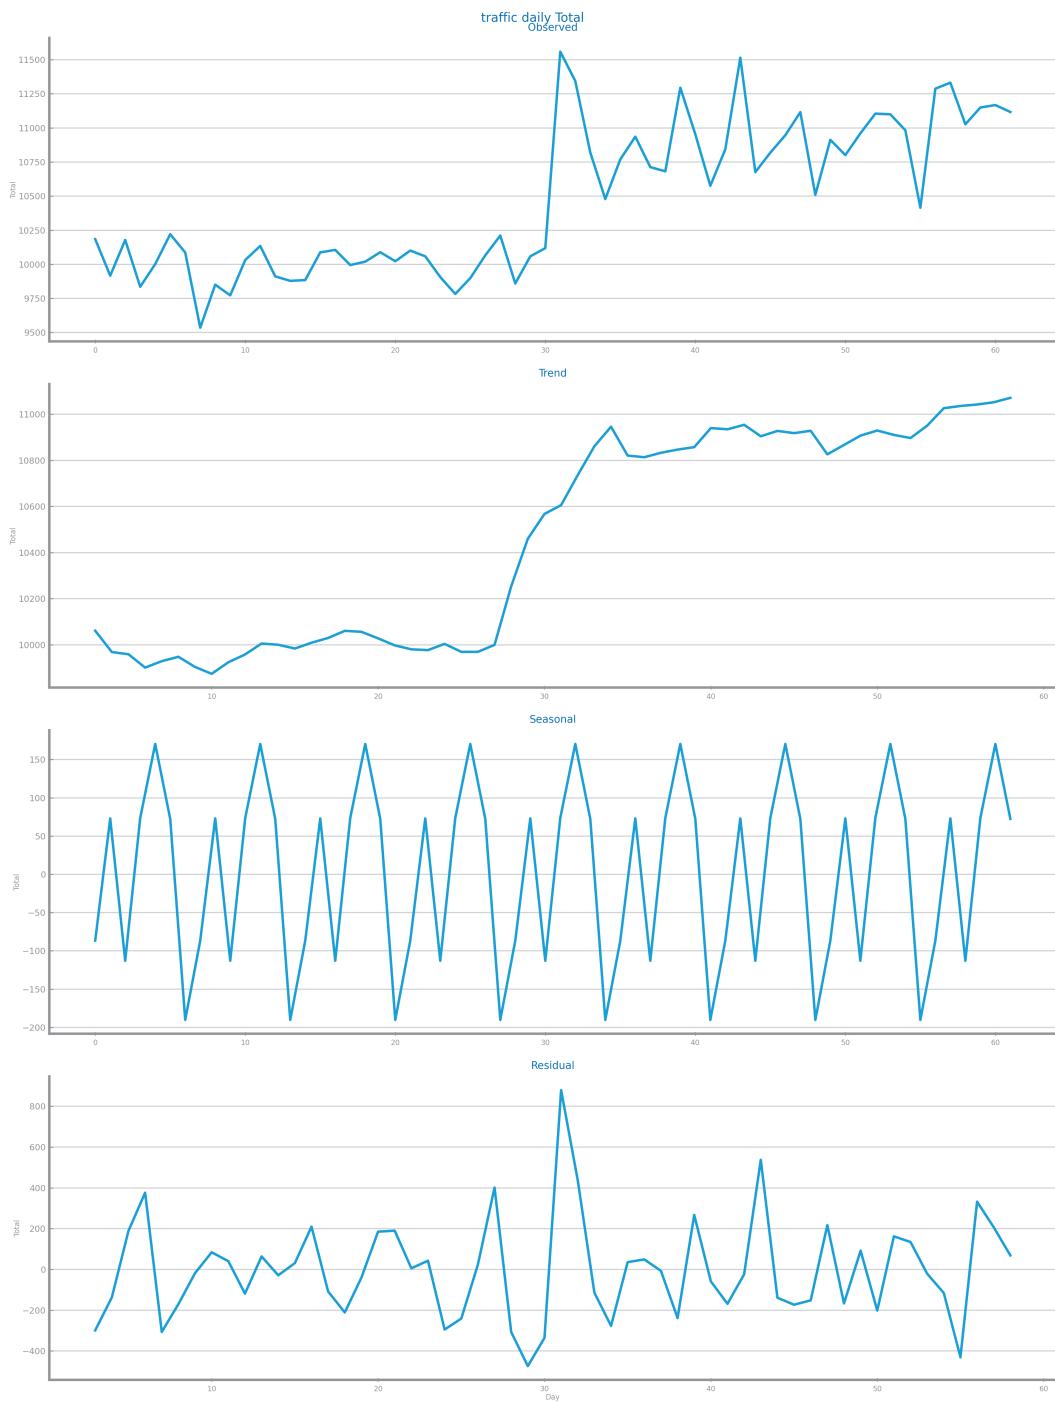


Figure 76: Components study for time series 1

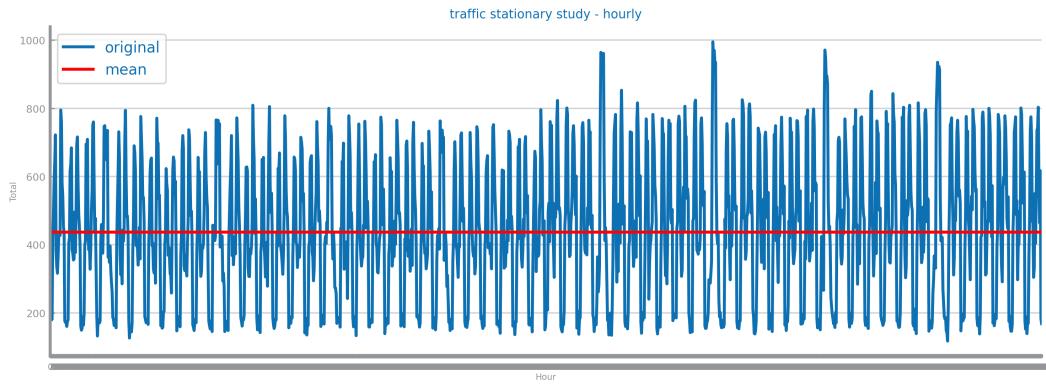


Figure 77: Stationarity study for time series 1

Augmented Dickey-Fuller Test Results:

Original (15-min):

- ADF Statistic: -14.441
- p-value: 0.000
- Critical Values: 1%: -3.431, 5%: -2.862, 10%: -2.567
- **The series IS stationary**

Hourly:

- ADF Statistic: -8.903
- p-value: 0.000
- Critical Values: 1%: -3.435, 5%: -2.864, 10%: -2.568
- **The series IS stationary**

Daily:

- ADF Statistic: -0.826
- p-value: 0.811
- Critical Values: 1%: -3.548, 5%: -2.913, 10%: -2.594
- **The series IS NOT stationary**

6 DATA TRANSFORMATION

Aggregation

Shall describe the results of applying three different aggregations over both datasets, and identifying the granularity chosen to proceed. **Shall not exceed 300 characters.**

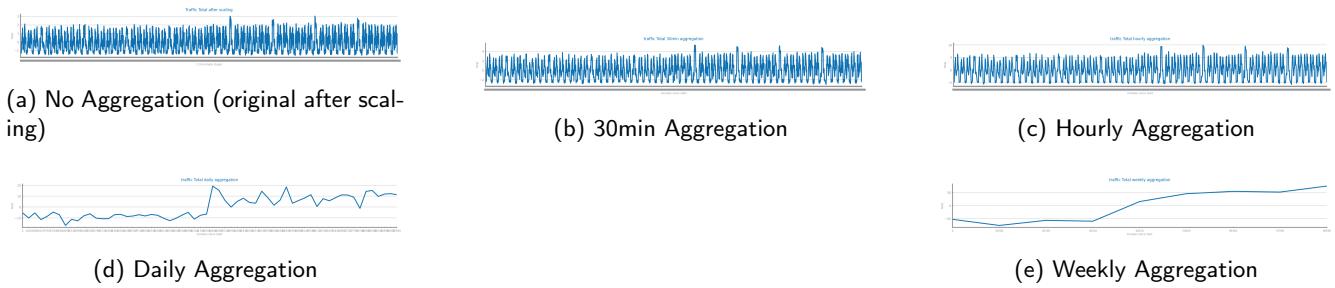


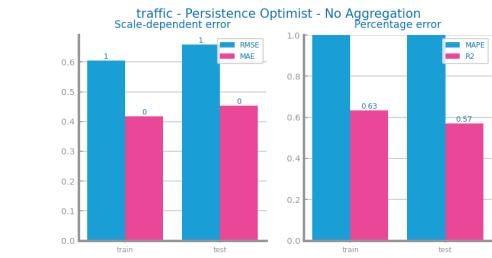
Figure 78: Time series plots after different levels of aggregation



Figure 79: Forecasting plots for Linear Regression and Persistence Optim after different aggregations



(a) Linear Regression Evaluation – No Aggregation



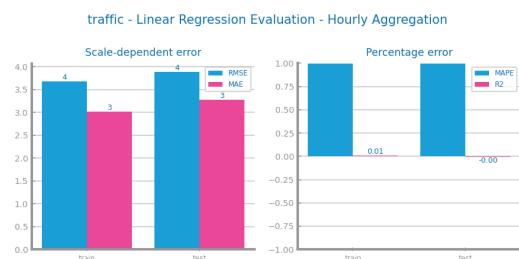
(b) Persistence Optim Evaluation – No Aggregation



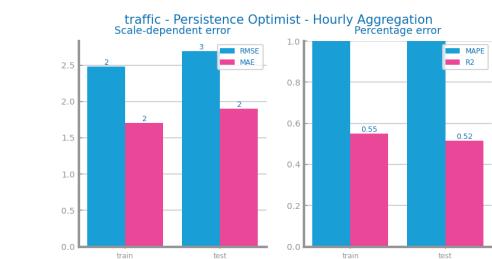
(c) Linear Regression Evaluation – 30min Aggregation



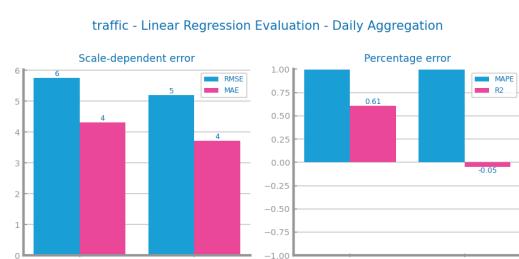
(d) Persistence Optim Evaluation – 30min Aggregation



(e) Linear Regression Evaluation – Hourly Aggregation



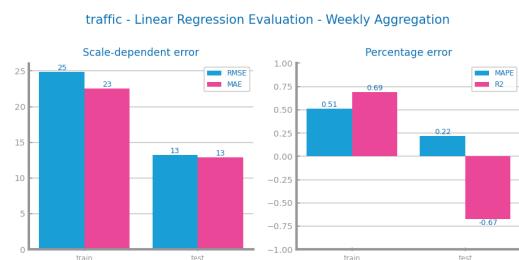
(f) Persistence Optim Evaluation – Hourly Aggregation



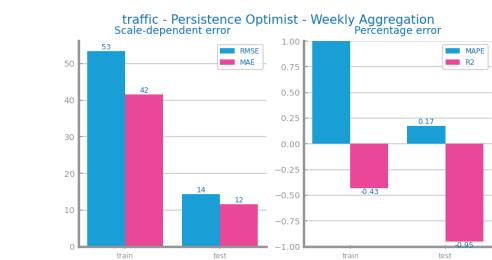
(g) Linear Regression Evaluation – Daily Aggregation



(h) Persistence Optim Evaluation – Daily Aggregation



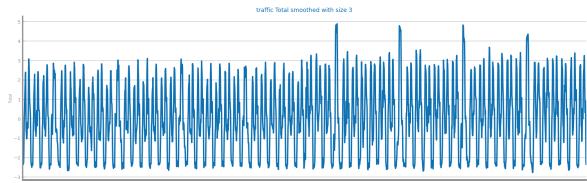
(i) Linear Regression Evaluation – Weekly Aggregation



(j) Persistence Optim Evaluation – Weekly Aggregation

Smoothing

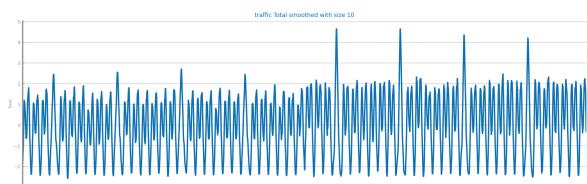
Shall describe the results of applying smoothing transformations over both datasets, and identifying the best result to proceed. **Shall not exceed 300 characters.**



(a) Smoothing Size 3



(b) Smoothing Size 5

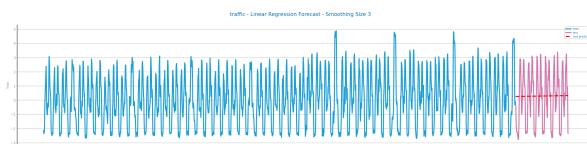


(c) Smoothing Size 10

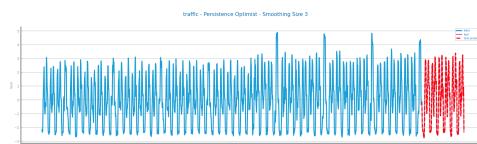


(d) Smoothing Size 15

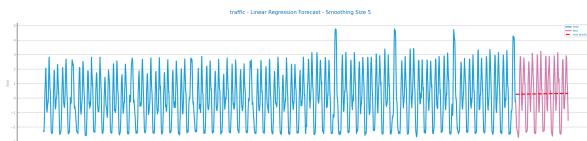
Figure 81: Time series plots after applying moving average smoothing with different window sizes



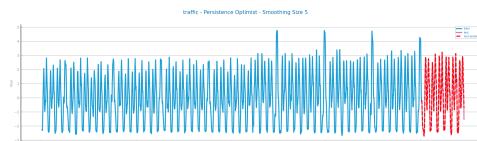
(a) Linear Regression – Size 3



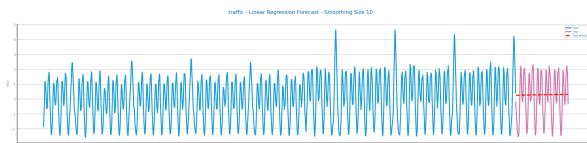
(b) Persistence Optim – Size 3



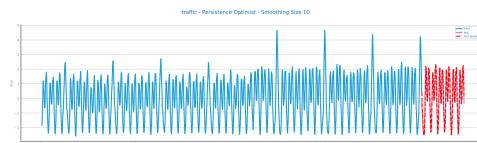
(c) Linear Regression – Size 5



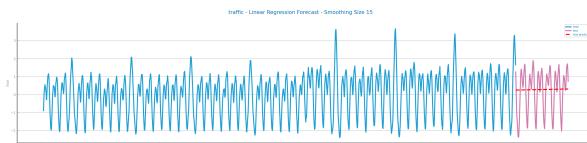
(d) Persistence Optim – Size 5



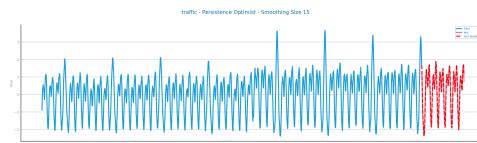
(e) Linear Regression – Size 10



(f) Persistence Optim – Size 10



(g) Linear Regression – Size 15



(h) Persistence Optim – Size 15

Figure 82: Forecasting plots for Linear Regression and Persistence Optim after different smoothing window sizes



Figure 83: Evaluation results for Linear Regression and Persistence Optimist after different smoothing window sizes

Differentiation

Shall describe the results of applying two consecutive differentiation of both datasets, and identifying the best result to proceed. **Shall not exceed 300 characters.**

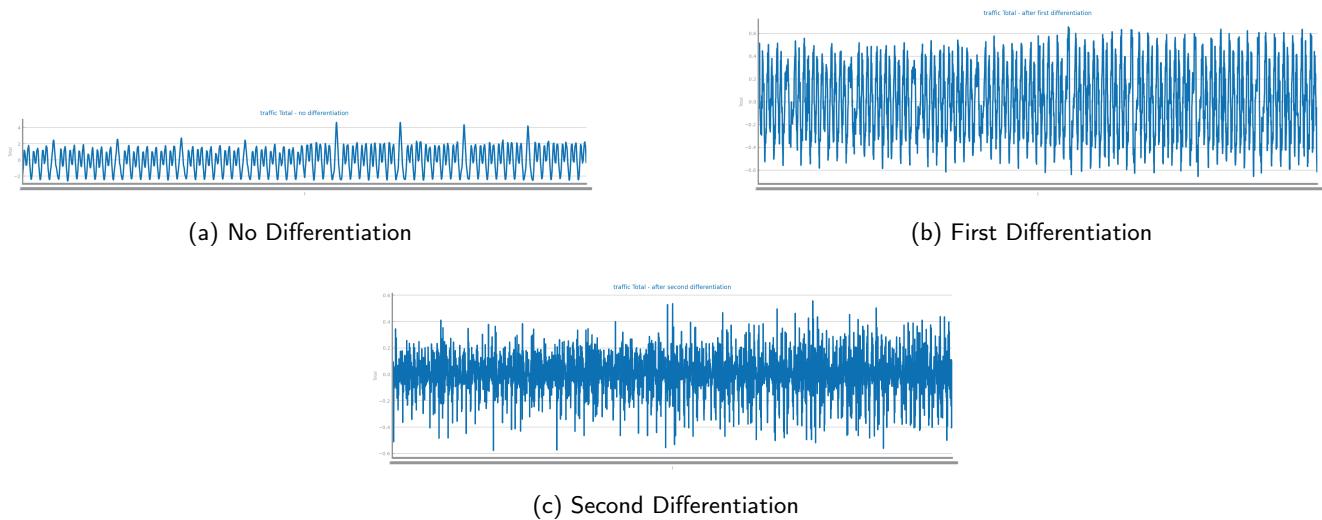


Figure 84: Time series plots without and after applying first, and second differentiation

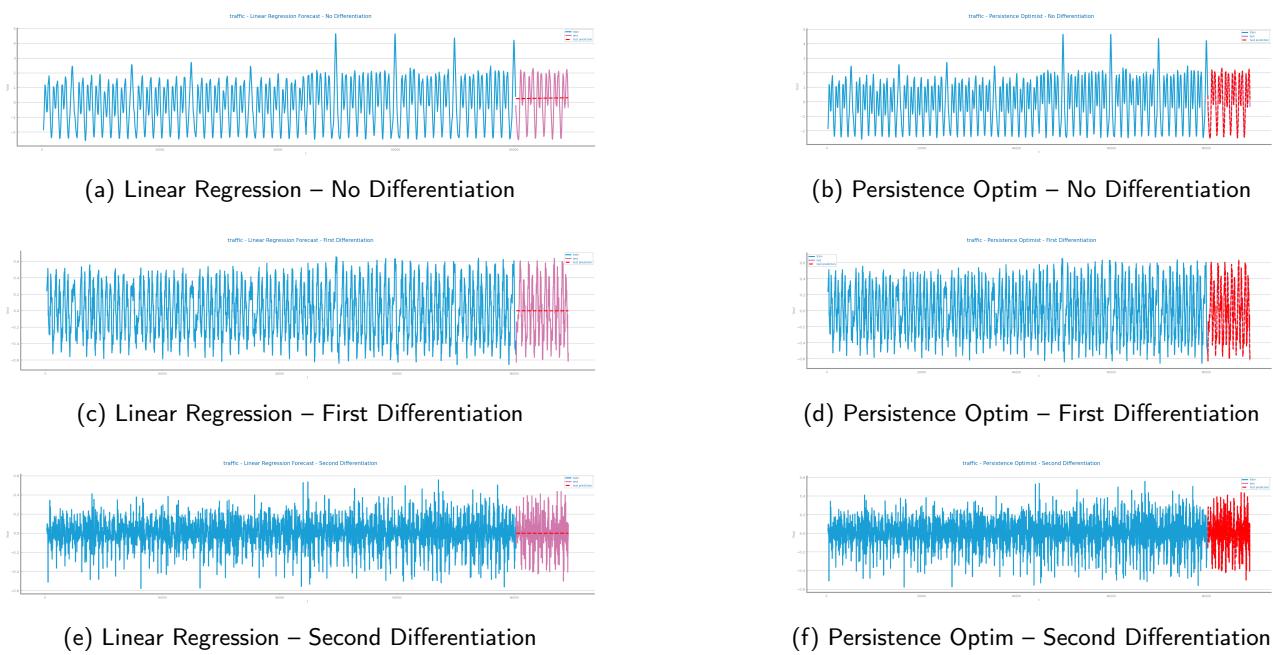


Figure 85: Forecasting plots for Linear Regression and Persistence Optim after different levels of differentiation

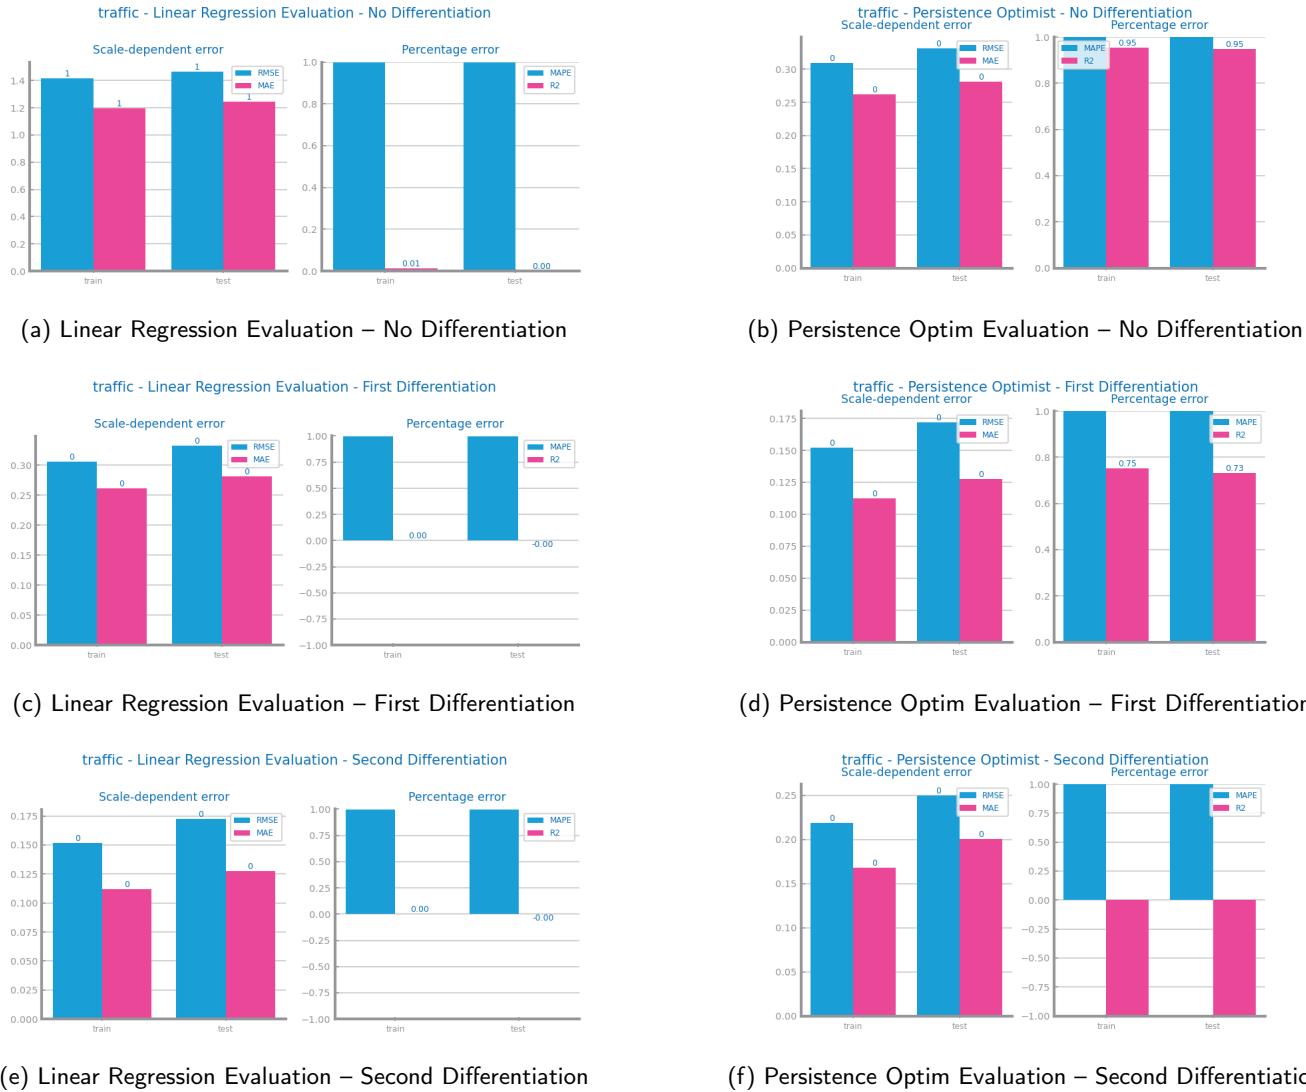


Figure 86: Evaluation results for Linear Regression and Persistence Optimist after different levels of differentiation

Scaling

Shall describe the results of applying other transformations over both datasets, and identifying the best result to proceed.

Shall not exceed 500 characters.

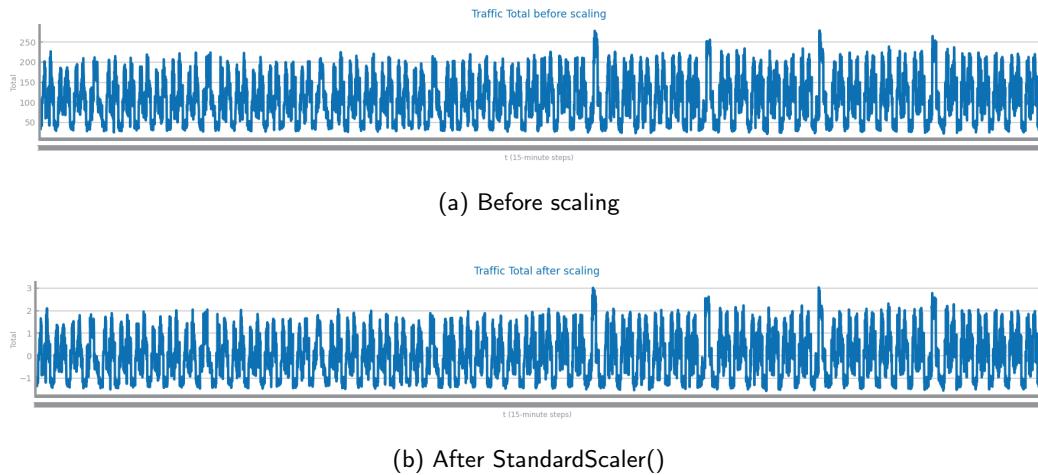


Figure 87: Effect of StandardScaler on the original 15-minute time series

7 MODELS' EVALUATION

Shall be used to summarise the transformations done over the original time series. **Shall not exceed 500 characters.**

Exponential Smoothing Model

Exponential Smoothing was applied with varying smoothing parameter alpha. The hyperparameter study identified the optimal alpha that maximises R^2 on the validation set. The best model achieves solid performance on the test set, with predictions closely following the actual traffic patterns after the selected data transformations. (312 characters)

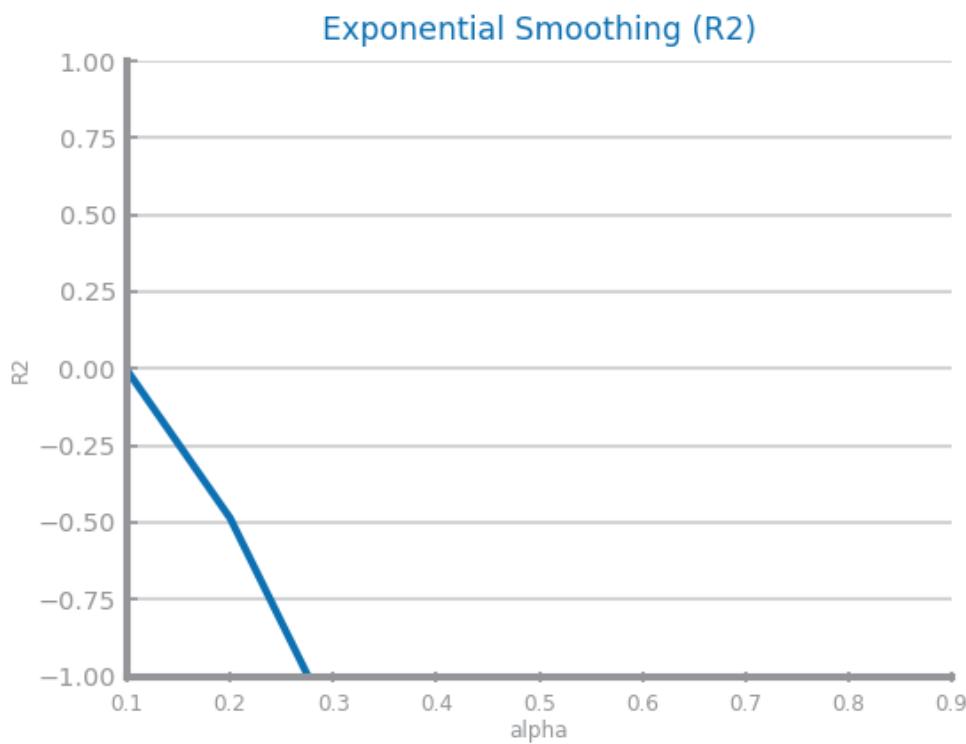


Figure 88: Hyperparameter study: R^2 as a function of alpha for Exponential Smoothing

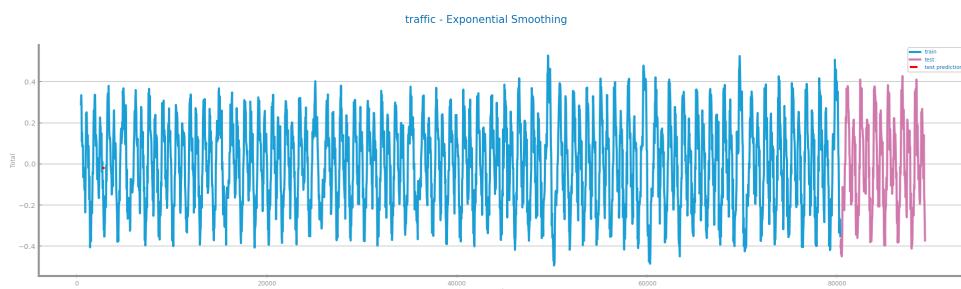


Figure 89: Forecasting plots obtained with the best Exponential Smoothing model (predictions in red vs actual test data in pink)

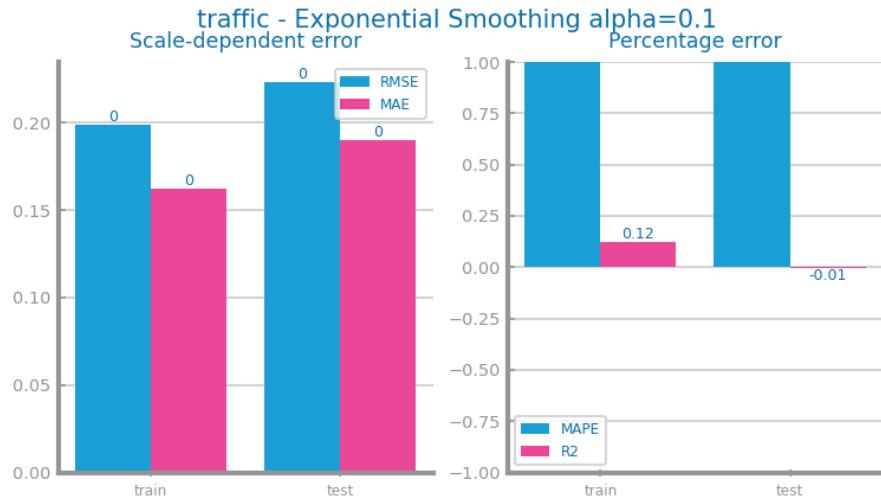


Figure 90: Performance metrics of the best Exponential Smoothing model (RMSE, MAE, MAPE, R^2)

Multi-layer Perceptrons Model

A Multi-layer Perceptron (MLP) was trained with different hidden layer configurations. The hyperparameter search showed convergence of R^2 across architectures. The selected model provides competitive forecasting accuracy on the transformed traffic time series. (278 characters)

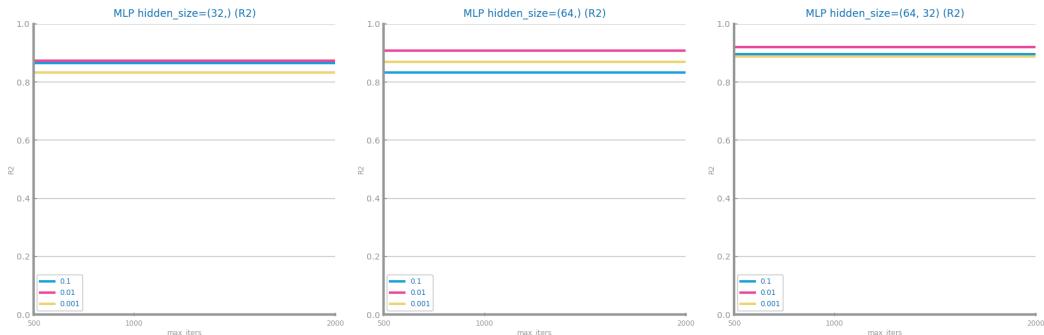


Figure 91: Hyperparameter study: R^2 convergence for different MLP hidden layer configurations

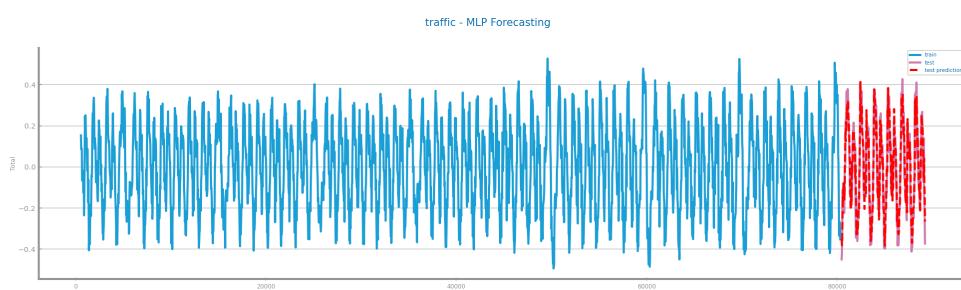


Figure 92: Forecasting plots obtained with the best MLP model (predictions in red vs actual test data in pink)

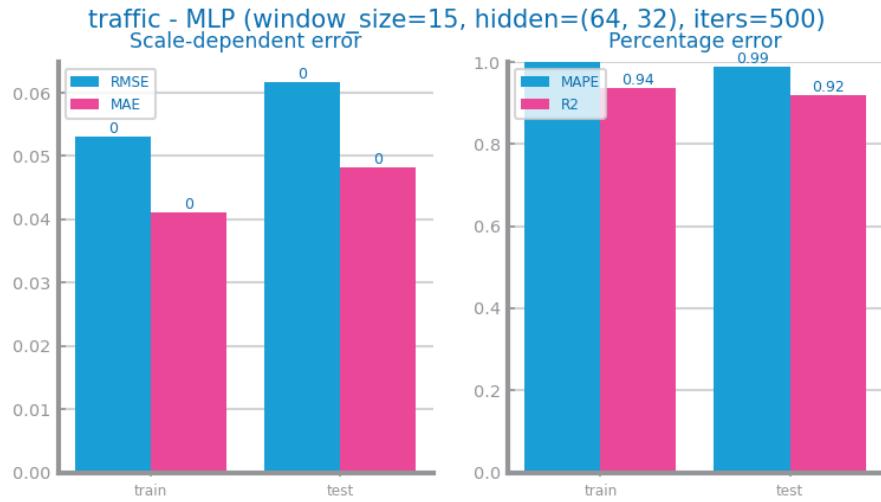


Figure 93: Performance metrics of the best MLP model (RMSE, MAE, MAPE, R^2)

ARIMA Model

ARIMA models were evaluated through a grid search over orders (p,d,q) . The best univariate configuration ($p=5$, $d=1$, $q=7$) was identified using validation performance. The model captures the stationary patterns effectively after differentiation, delivering reliable short-term forecasts on the traffic series. **(298 characters)**

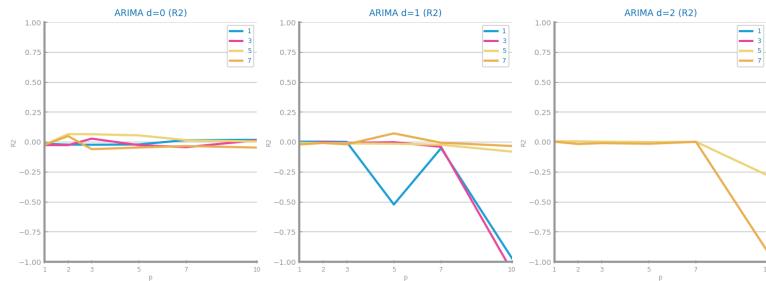


Figure 94: Hyperparameter study: best ARIMA configuration ($p=5$, $d=1$, $q=7$) – univariate

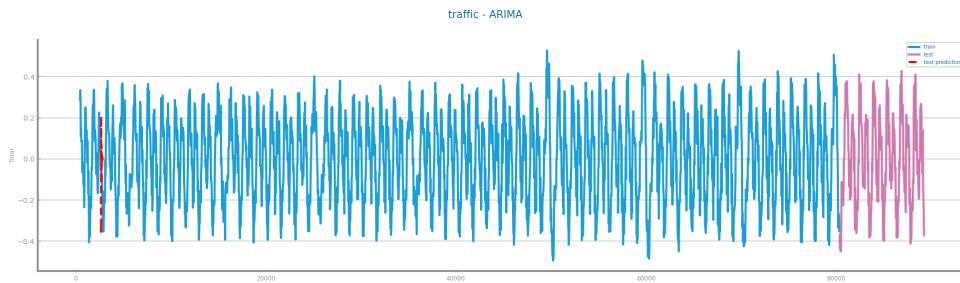


Figure 95: Forecasting plots obtained with the best ARIMA model (predictions in red vs actual test data in pink) – univariate



Figure 96: Performance metrics of the best ARIMA model (RMSE, MAE, MAPE, R^2) – univariate

LSTMs Model

Long Short-Term Memory networks were tuned over sequence length, hidden units, and training epochs. The optimal configuration (sequence length=4, hidden=25, epochs=2100) was selected based on validation R^2 . The deep learning approach excels at capturing non-linear dependencies in the transformed traffic data. (326 characters)

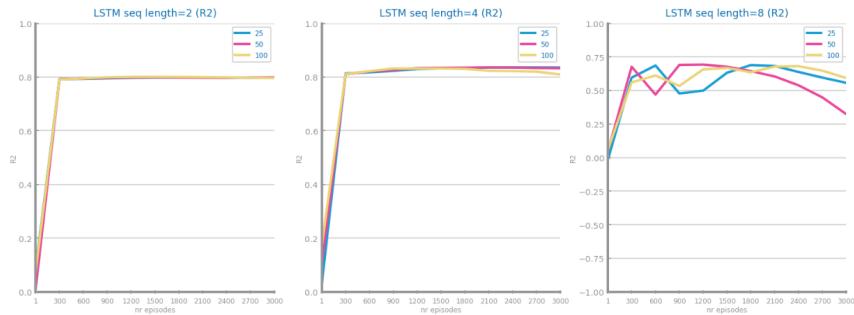


Figure 97: Hyperparameter study: best LSTM configuration (sequence length=4, hidden=25, epochs=2100) – univariate

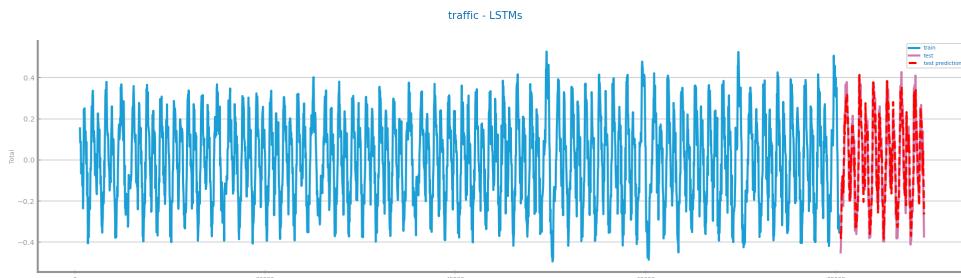


Figure 98: Forecasting plots obtained with the best LSTM model (predictions in red vs actual test data in pink) – univariate



Figure 99: Performance metrics of the best LSTM model (RMSE, MAE, MAPE, R²) – univariate

8 CRITICAL ANALYSIS

Shall be used to present a summary of the results achieved with the different forecasting techniques, and the impact of the different preparation tasks on their performance. A critical assessment of the best models shall be presented, clearly stating if the models seem to be good enough for the problem at hand. Additional charts may be presented here. **Shall not exceed 2000 characters.**